CLAUDE CERTIFIED ARCHITECT - FOUNDATIONS | EXAM CODE CCAR-F

Mock Examination 05

Claude Certified Architect - Foundations
60 scenario-based questions | 120 minutes

Scenarios assessed in this paper: Scenario 2 - Code Generation with Claude Code; Scenario 3 - Multi-Agent Research System; Scenario 5 - Claude Code for Continuous Integration; Scenario 6 - Structured Data Extraction

Format: multiple-choice and multiple-response; every question states how many answers to select.

Contents: instructions and blueprint | candidate answer sheet | 60 questions | quick answer key and score tracker | detailed explanations for every option | domain-level review guide | revision checklist.

Status: ORIGINAL PRACTICE MATERIAL. This is not an official Anthropic examination paper and contains no confidential exam content. It was written against the publicly available Claude Certified Architect - Foundations Exam Guide (v1.0, effective July 2026).

Contents

1. Instructions and exam conditions 3

2. What is verified, what is supplied, and what is original 3

3. Exam blueprint for this paper 4

Candidate answer sheet 5

4. Examination - Questions 1 to 60 6

Scenario 2: Code Generation with Claude Code (Questions 1-13) . . . 6

Scenario 3: Multi-Agent Research System (Questions 14-29) . . . 8

Scenario 5: Claude Code for Continuous Integration (Questions 30-43) . . . 12

Scenario 6: Structured Data Extraction (Questions 44-60) . . . 16

5. Quick answer key and score tracker 20

6. Detailed answer key and explanations 21

7. Domain-level performance review guide 45

Domain 1 - Agentic Architecture & Orchestration (27% official weight; 16 questions here) . . . 45

Domain 2 - Tool Design & MCP Integration (18% official weight; 11 questions here) . . . 45

Domain 3 - Claude Code Configuration & Workflows (20% official weight; 12 questions here) . . . 45

Domain 4 - Prompt Engineering & Structured Output (20% official weight; 12 questions here) . . . 45

Domain 5 - Context Management & Reliability (15% official weight; 9 questions here) . . . 46

8. Final revision checklist 47

1. Instructions and exam conditions

Time limit: 120 minutes for 60 questions (an average of 2 minutes per question). Work in one sitting, without notes, documentation, or AI assistance, to get a realistic read on your readiness.

• Every question is anchored to one of the production scenarios in this paper. Read the scenario description at the start of each scenario section; questions are otherwise self-contained.
• Each question states **how many answers to select**. For 'Select ONE', exactly one option is best. For 'Select exactly TWO/THREE', all selected options must be correct; there is no partial credit in this practice format.
• Choose the *best* answer for the stated constraints, not an answer that could merely work. Where a question asks for the FIRST, SAFEST, or MOST RELIABLE step, that qualifier is part of the question.
• Attempt every question. Unanswered questions score zero. There is no 'all of the above' or 'none of the above'.
• Do not look at the answer key until you have finished. The key, explanations, and review guide begin after Question 60.

2. What is verified, what is supplied, and what is original

Category	What it covers
Verified official information	From the official Exam Guide (CCAR-F v1.0): 60 items; multiple-choice and multiple-response items that state how many responses to select; 4 scenarios drawn from a bank of 6; 120-minute limit; scaled passing score of 720 on a 100-1,000 scale; score report with percent-correct by domain; domain weights 27/18/20/20/15%; the 31 task statements and the in-scope/out-of-scope lists.
Scenario descriptions	The six production scenarios are quoted from the official guide; the extra topic lists you were given (e.g., idempotency, structured handoffs, schema design) guided question themes. Scenario text shown above each question group is the official description.
Original practice content	All questions, options, explanations, and review guidance in this file were written for this pack. They are not official questions, are not copied from commercial practice exams or leaked banks, and may differ in difficulty from the real exam.
Supplemental knowledge	A few items use general engineering knowledge to make a scenario realistic (for example CI/CD terminology or least-privilege shell access). They are flagged 'Supplemental' in the explanation and are not confirmed exam objectives.

A raw percentage on this paper is **not** a scaled score. The real exam reports a scaled score (100-1,000) with a cut score of 720 derived from a formal standard-setting study; percent-correct by domain is shown for information only. Doing well on these papers does not guarantee a pass.

3. Exam blueprint for this paper

Domain	Official weight	Questions in this paper
1. Agentic Architecture & Orchestration	27%	16
2. Tool Design & MCP Integration	18%	11
3. Claude Code Configuration & Workflows	20%	12
4. Prompt Engineering & Structured Output	20%	12
5. Context Management & Reliability	15%	9
Total	**100%**	**60**

Question counts are the closest integer allocation of the official weights to 60 questions. Each question has exactly one primary domain; secondary domains appear in the answer key where a question also touches another domain.

Scenario	Questions	Numbers
Scenario 2 - Code Generation with Claude Code	13	1-13
Scenario 3 - Multi-Agent Research System	16	14-29
Scenario 5 - Claude Code for Continuous Integration	14	30-43
Scenario 6 - Structured Data Extraction	17	44-60

Design targets (not claims about the real exam): about 20% foundational, 55% intermediate, 25% challenging. In this paper: 12 foundational, 33 intermediate, 15 challenging. Response types: 54 single-answer, 5 select-two, 1 select-three.

Candidate answer sheet

Mark your answers here while taking the paper. Cells that do not exist for a question are greyed out. Where a question says 'Select exactly TWO/THREE', mark exactly that many letters. Do not turn to the answer key until you have finished all 60 questions (or the 120 minutes have elapsed).

Name: _____ Date: _____ Start time: _____ End time: _____

No.	Pick	A	B	C	D	E	F
1	1	(A)	(B)	(C)	(D)		
2	1	(A)	(B)	(C)	(D)		
3	1	(A)	(B)	(C)	(D)		
4	1	(A)	(B)	(C)	(D)		
5	1	(A)	(B)	(C)	(D)		
6	2	(A)	(B)	(C)	(D)	(E)	
7	2	(A)	(B)	(C)	(D)	(E)	
8	2	(A)	(B)	(C)	(D)	(E)	
9	1	(A)	(B)	(C)	(D)		
10	1	(A)	(B)	(C)	(D)		
11	1	(A)	(B)	(C)	(D)		
12	1	(A)	(B)	(C)	(D)		
13	1	(A)	(B)	(C)	(D)		
14	1	(A)	(B)	(C)	(D)		
15	1	(A)	(B)	(C)	(D)		
16	1	(A)	(B)	(C)	(D)		
17	1	(A)	(B)	(C)	(D)		
18	1	(A)	(B)	(C)	(D)		
19	1	(A)	(B)	(C)	(D)		
20	1	(A)	(B)	(C)	(D)		
21	1	(A)	(B)	(C)	(D)		
22	1	(A)	(B)	(C)	(D)		
23	1	(A)	(B)	(C)	(D)		
24	1	(A)	(B)	(C)	(D)		
25	2	(A)	(B)	(C)	(D)	(E)	
26	1	(A)	(B)	(C)	(D)		
27	3	(A)	(B)	(C)	(D)	(E)	(F)
28	1	(A)	(B)	(C)	(D)		
29	1	(A)	(B)	(C)	(D)		
30	1	(A)	(B)	(C)	(D)		

No.	Pick	A	B	C	D	E	F
31	1	(A)	(B)	(C)	(D)		
32	1	(A)	(B)	(C)	(D)		
33	1	(A)	(B)	(C)	(D)		
34	1	(A)	(B)	(C)	(D)		
35	1	(A)	(B)	(C)	(D)		
36	1	(A)	(B)	(C)	(D)		
37	1	(A)	(B)	(C)	(D)		
38	1	(A)	(B)	(C)	(D)		
39	1	(A)	(B)	(C)	(D)		
40	1	(A)	(B)	(C)	(D)		
41	1	(A)	(B)	(C)	(D)		
42	1	(A)	(B)	(C)	(D)		
43	1	(A)	(B)	(C)	(D)		
44	1	(A)	(B)	(C)	(D)		
45	1	(A)	(B)	(C)	(D)		
46	1	(A)	(B)	(C)	(D)		
47	1	(A)	(B)	(C)	(D)		
48	1	(A)	(B)	(C)	(D)		
49	1	(A)	(B)	(C)	(D)		
50	1	(A)	(B)	(C)	(D)		
51	1	(A)	(B)	(C)	(D)		
52	1	(A)	(B)	(C)	(D)		
53	1	(A)	(B)	(C)	(D)		
54	1	(A)	(B)	(C)	(D)		
55	1	(A)	(B)	(C)	(D)		
56	1	(A)	(B)	(C)	(D)		
57	2	(A)	(B)	(C)	(D)	(E)	
58	1	(A)	(B)	(C)	(D)		
59	1	(A)	(B)	(C)	(D)		
60	1	(A)	(B)	(C)	(D)		

4. Examination - Questions 1 to 60

Time allowed: 120 minutes. Start the clock now.

Scenario 2: Code Generation with Claude Code (Questions 1-13)

You are using Claude Code to accelerate software development. Your team uses it for code generation, refactoring, debugging, and documentation. You need to integrate it into your development workflow with custom slash commands, CLAUDE.md configurations, and understand when to use plan mode vs direct execution.

Question 1

Select ONE answer.

A Claude Code investigation of flaky tests ran for several days. Meanwhile, the test runner was upgraded and many earlier tool outputs now describe behavior that no longer exists. You decide to start a new session with a structured summary. Which content should the summary carry?

A. Only the first and last message of the old session, since they contain the goal and the final result.

B. Every command the old session ran, in order, so that the new session can replay them and compare results.

C. The complete transcript of the old session, pasted in full, so that nothing the previous session saw is lost and the new session has every detail available.

D. Confirmed conclusions, decisions already made, files of interest, and open questions, leaving out raw tool output tied to the old runner.

Question 2

Select ONE answer.

Which situation is the best use of fork_session?

A. Comparing two candidate refactoring approaches, starting from the same completed analysis of a module.

B. Recovering from a crash by restoring the last saved state of every tool call.

C. Continuing yesterday's debugging work on the same bug with the same accumulated context.

D. Beginning an unrelated task in a different repository with no connection to the earlier analysis.

Question 3

Select ONE answer.

A team lists four pieces of guidance: (a) a developer's preference for terse replies, (b) rules that apply only to files under payments/, (c) naming standards for the entire repository, (d) Terraform module conventions for .tf files spread across several folders. Which mapping to Claude Code mechanisms is most appropriate?

A. (a) a skill, (b) the project-level CLAUDE.md, (c) the user-level ~/.claude/CLAUDE.md, (d) a directory-level CLAUDE.md in terraform/.

B. (a) directory-level CLAUDE.md in payments/, (b) the project-level CLAUDE.md, (c) a .claude/rules/ file without paths, (d) the user-level ~/.claude/CLAUDE.md.

C. (a) user-level ~/.claude/CLAUDE.md, (b) directory-level CLAUDE.md in payments/, (c) project-level CLAUDE.md, (d) a .claude/rules/ file scoped by a **/*.tf glob.

D. (a) project-level CLAUDE.md, (b) user-level ~/.claude/CLAUDE.md, (c) directory-level CLAUDE.md in each package, (d) a skill invoked manually before editing any Terraform file.

Question 4

Select ONE answer.

A teammate asks which Claude Code feature supports frontmatter options like context: fork, allowed-tools, and argument-hint. What is the answer?

A. The settings file for hooks, as fields on the PostToolUse entry.

B. Skills, defined as SKILL.md files under .claude/skills/.

C. The project-level CLAUDE.md file, in a block at the top of the file.

D. The .mcp.json file, as properties on each server entry.

Question 5

Select ONE answer.

A developer is facing three problems: (i) Claude interprets a prose description of a data transformation inconsistently, (ii) the developer is implementing something in a domain they don't know well, (iii) a function mishandles edge cases and needs a verifiable standard. Which pairing of problem and technique is best?

A. (i) write tests first and iterate on failures, (ii) provide concrete input/output examples, (iii) use the interview pattern.

B. (i) provide two or three concrete input/output examples, (ii) use the interview pattern, (iii) write tests first and iterate on failures.

C. (i) use the interview pattern, (ii) write tests first and iterate on failures, (iii) provide two or three concrete input/output examples for each edge case.

D. (i) repeat the prose with greater emphasis, (ii) provide one long example, (iii) ask Claude to rate its confidence.

Question 6

Select exactly TWO answers (choose 2).

A new engineer asks how path-specific rule files in .claude/rules/ behave in practice. Select the TWO accurate statements.

A. Path-scoped rules load only when Claude works on matching files, which reduces irrelevant context and token usage.

B. A paths value can only be an exact file name, so a glob pattern isn't allowed.

C. A rule with paths: ["**/*.test.*"] applies to test files wherever they live in the repository.

D. A directory-level CLAUDE.md automatically applies to matching files in sibling directories.

E. Path-scoped rules only apply after a developer runs a slash command naming the rule file.

Question 7

Select exactly TWO answers (choose 2).

Select the TWO tasks that are best suited to direct execution rather than plan mode.

A. Adding a date validation conditional to one existing function with fully specified behavior.

B. Choosing between two integration approaches that need different infrastructure.

C. Restructuring a monolith into microservices, which requires decisions about service boundaries.

D. Migrating a logging library across roughly 45 files with several custom wrappers.

E. Fixing a crash where the stack trace points at one specific line in one function.

Question 8

Select exactly TWO answers (choose 2).

A team is documenting how skill frontmatter options work and where skill and command files live. Select the TWO accurate statements.

A. allowed-tools in the frontmatter restricts which tools the skill may use while it runs.

B. Commands in ~/.claude/commands/ are shared with teammates through version control.

C. Skills in ~/.claude/skills/ replace the project's skill of the same name for the whole team.

D. argument-hint restricts which tools a skill is allowed to use.

E. context: fork runs a skill in an isolated sub-agent context so verbose output doesn't pollute the main conversation.

Question 9

Select ONE answer.

A rule file has the frontmatter paths: ["terraform/**/*"]. When will Claude Code load this rule?

A. When it is working with files that match terraform/**/*, such as an edit to terraform/network/main.tf.

B. At the start of every session, regardless of which files are being edited.

C. Only for files directly inside terraform/, not for files in subdirectories.

D. Only when the developer types the rule's file name as a slash command.

Question 10

Select ONE answer.

A function must process 10,000 items within 200 ms and handle an empty list, a single item, and duplicate items. Previous attempts passed the happy path but were slow. Which workflow gives the most targeted feedback loop?

A. Describe the performance requirement in more detail and ask Claude to reason carefully about algorithmic complexity before writing any code, then ship it.

B. Implement first, then ask Claude to produce tests that pass for the implementation.

C. Write tests for correctness, edge cases, and the 200 ms requirement first, then have Claude implement and iterate by sharing failing test output.

D. Ask for three different implementations and choose the one that looks the most efficient.

Question 11

Select ONE answer.

Claude Code must decide whether a new behavior calls for a unit test or an integration test. The choice is often ambiguous, and results vary. Which prompt addition is most helpful?

A. A reminder to think carefully before choosing and to be consistent from file to file.

B. A few examples of ambiguous cases showing which test type was chosen and why it was preferred over the alternative.

C. A rule that every behavior gets both a unit test and an integration test to avoid choosing.

D. An instruction to pick whichever test type is faster to write in each case.

Question 12

Select ONE answer.

The team's few-shot examples for Claude Code were copied from old modules that use callbacks. New generated code also uses callbacks, although the current house style is async/await. What is the right fix?

A. Add more callback examples to improve consistency, then rely on a code formatter to rewrite them.

B. Keep the callback examples and add a sentence to the prompt saying that async/await is the preferred style for all new code written from now on.

C. Replace the examples with ones drawn from current, preferred code written in async/await, since examples strongly shape generated style.

D. Remove all examples so that the model uses its own default style instead.

Question 13

Select ONE answer.

A developer asks what /compact does in a long Claude Code session. Which description is correct?

A. It reduces context usage by condensing the conversation so a long session can continue.

B. It permanently deletes the project's CLAUDE.md to save space in the repository.

C. It compresses source files on disk so that they take less space.

D. It restarts Claude Code with a different model that has a larger context window.

Scenario 3: Multi-Agent Research System (Questions 14-29)

You are building a multi-agent research system using the Claude Agent SDK. A coordinator agent delegates to specialized subagents: one searches the web, one analyzes documents, one synthesizes findings, and one generates reports. The system researches topics and produces comprehensive, cited reports.

Question 14

Select ONE answer.

In an agentic loop for a research subagent, what should the code do when the response's stop_reason is 'tool_use'?

A. Discard the tool request and call the model again with the same messages to get a different response.

B. Execute the requested tool(s), append the tool results to the conversation, and send another request.

C. Return the response text to the coordinator as the final answer, since the model has finished.

D. Stop the loop and log an error because the subagent should not use tools.

Question 15

Select ONE answer.

A research agent normally ends via stop_reason 'end_turn', but a few runs wander through dozens of searches. The team wants cost protection without cutting off legitimate deep research silently. What is the best design?

A. Use a hard iteration cap of 5 as the primary stopping rule, and return whatever has been gathered by then as the final report for the user without comment.

B. Stop the loop whenever the model writes a sentence beginning 'In summary', since that signals completion.

C. Keep stop_reason as the normal terminator, and add token and tool-call budgets that stop gracefully and produce a partial report stating what was not covered.

D. Remove all limits and rely on the model to decide when it has searched enough.

Question 16

Select ONE answer.

A coordinator spawns the synthesis subagent. Which of the following is the most complete context packet to include in its prompt?

A. A short summary of the findings in prose, without sources, to keep the prompt compact.

B. The research question and quality criteria, the findings as structured records (claim, excerpt, source, date), noted conflicts, and the required format.

C. The research question and the instruction 'summarize everything found so far', with no findings included in the prompt.

D. The complete raw text of all fetched pages, with no structure, so that no detail is lost.

Question 17

Select ONE answer.

Two web research subagents were asked to research 'AI regulation'. Their results overlap heavily, leaving other relevant areas thin. How should the coordinator change its delegation?

A. Keep the same assignment but have the synthesis subagent remove duplicate results later.

B. Replace the two subagents with a single one that researches the whole topic sequentially.

C. Ask both subagents to work harder and to be more thorough, in the hope that they will find different material by chance on the next run, without changing assignments.

D. Partition the scope: assign distinct subtopics or source types to each subagent, such as regions, regulatory bodies, and industry commentary.

Question 18

Select ONE answer.

A pre-call hook blocks the web fetch tool for a disallowed domain by raising an exception that ends the agent loop. The user receives an error instead of a report. What is the better behavior?

A. Remove the hook, since blocking domains should be handled by telling the model in the prompt.

B. Silently return an empty result for the blocked call so the model assumes there is nothing at that domain.

C. Let the exception end the loop, because a blocked call should always abort the whole task for safety reasons and report the policy violation to the user.

D. Have the hook return a tool_result explaining that the call was blocked and why, so the model can continue with another source or note the gap.

Question 19

Select ONE answer.

The research workflow has these steps: web search, document analysis of the user's PDFs, synthesis, and report writing. Web search and document analysis are independent of each other. Which execution plan is most appropriate?

A. Run all four steps in parallel and merge the results at the end.

B. Run all four steps strictly one after another, so every step sees the output of the one before.

C. Run web search and document analysis in parallel, then run synthesis once both complete, then report writing.

D. Run synthesis first so that it can tell search and analysis what to look for.

Question 20

Select ONE answer.

The coordinator invokes the document analysis subagent for batch 1 of PDFs, and later invokes it again for batch 2, assuming the subagent will remember the entities identified in batch 1. The batch 2 results contradict batch 1 naming. What is the correct fix?

A. Run batch 2 before batch 1 so that the more complex documents are processed first.

B. Tell the subagent in the batch 2 prompt to remember what it learned during batch 1 and to keep the entity names consistent with its earlier work, whatever they were.

C. Use a larger model for batch 2 because it can better infer earlier decisions.

D. Pass the batch 1 results (entity names and mappings) explicitly in the batch 2 prompt, since subagents don't share memory between invocations.

Question 21

Select ONE answer.

A research session on regional energy prices was started three days ago. The searches returned news and price tables that have since been updated. The user asks for a refreshed report. Which approach is most reliable?

A. Resume the old session and tell the model that some prices may have changed, then use its best judgment.

B. Fork the old session and ask the fork to trust the earlier tool results.

C. Resume the old session and ask for a refreshed report using the data that is already in its history, which saves the cost of repeating the searches.

D. Start a new session with a structured summary of the research plan and verified background, and re-run the searches for time-sensitive data.

Question 22

Select ONE answer.

The web search subagent requests a research article and receives HTTP 403 from a paywalled site. The tool returns 'Request failed'. The agent retries three times. What is the best tool-level response, and the right recovery?

A. Return an empty successful result so the subagent moves on without any message.

B. Return a transient error with isRetryable true so that the subagent keeps trying with increasing delays until the paywall eventually disappears or the access token is refreshed.

C. Return the HTML of the paywall page as the article content so the subagent can try to read it.

D. Return a permission-category, non-retryable error naming the cause; the subagent seeks an open-access version or another source, and reports a gap if none exists.

Question 23

Select ONE answer.

The team has an MCP tool search_licensed_papers that searches a licensed database and returns full metadata and abstracts. Its description reads 'Searches papers.' The research agent keeps using the built-in WebSearch instead. What should be done?

A. Remove WebSearch from all agents, so the licensed tool is the only search option anywhere in the system and no agent can be tempted by a built-in tool.

B. Add a rule in the system prompt: 'never use WebSearch under any circumstances'.

C. Expand the MCP tool's description to explain what it covers, what it returns, and when it is preferable to a general web search.

D. Rename the tool to something longer and hope that its name attracts the model.

Question 24

Select ONE answer.

The research coordinator has been given the web search, document loading, and verification tools of its subagents plus 20 other tools, 'for flexibility'. It increasingly does research itself instead of delegating and picks tools poorly. What is the best change?

A. Keep the tool set but add a paragraph to the prompt telling the coordinator not to use specialist tools directly and to delegate such work.

B. Give the subagents the same 20 extra tools so that everyone has the same capabilities.

C. Force tool_choice to the Task tool on every turn so that the coordinator always delegates.

D. Restrict the coordinator to a small set of orchestration tools, chiefly Task, and leave the specialist tools with the subagents.

Question 25

Select exactly TWO answers (choose 2).

Select the TWO elements that most improve a tool description for reliable selection among similar tools.

A. Input formats and example queries that illustrate typical use.

B. Enthusiastic wording stating that the tool is fast, powerful, and reliable.

C. The tool's internal ticket number and the date it was last deployed.

D. Boundary statements describing when to use this tool versus a similar one.

E. A reduction of the description to under five words for readability.

Question 26

Select ONE answer.

Two tools in the research agent, get_info and fetch_data, have nearly identical descriptions. The first actually queries an internal knowledge base and the second downloads public web pages. The agent uses them interchangeably. What is the best fix?

A. Combine the two into a single get_data tool with a source argument.

B. Rename them to reflect their roles (for example search_internal_kb and fetch_web_page) and rewrite each description with its source, inputs, outputs, and use.

C. Add a third tool that decides which of the two to call and then calls it.

D. Keep the names and add a one-line prompt: 'prefer get_info for questions about the company and fetch_data for questions about the outside world' whenever possible.

Question 27

Select exactly THREE answers (choose 3).

Select the THREE statements about review architectures that are correct.

A. A model that generated content retains its reasoning context, which makes it less likely to question its own decisions in the same session.

B. Splitting a large review into per-file passes plus a cross-file integration pass helps avoid attention dilution and contradictory findings.

C. A larger context window eliminates attention dilution, so a large review no longer needs to be split.

D. Telling the same session to act as a skeptical reviewer is equivalent to using a separate instance.

E. Extended thinking removes the need for an independent reviewer on important changes.

F. An independent instance without the generator's reasoning context is more effective at catching subtle issues than self-review instructions or extended thinking.

Question 28

Select ONE answer.

A final research report lists all findings in one uniform bulleted list. Some bullets come from multiple strong sources, while one topic area had no accessible sources at all, and two bullets are contested between sources. Readers cannot tell them apart. How should the synthesis be structured?

A. Remove contested findings and gaps from the report to keep it clean and authoritative.

B. Keep the single list but order the bullets by importance, trusting the reader to infer the level of confidence from each bullet's position in the list.

C. Use separate sections for well-established findings, contested findings with each source's position, and a coverage section stating gaps from unavailable sources.

D. Add a footer note saying 'some findings may be uncertain' without identifying which ones.

Question 29

Select ONE answer.

A subagent searched 12 reputable sources for studies on a narrow topic and found none. What should it report?

A. Nothing at all, since an empty result has no value to the coordinator.

B. A failure with category transient and isRetryable true, so that the coordinator retries the search repeatedly for this topic until something is found.

C. A fabricated summary of related studies to avoid reporting an empty result.

D. A successful result with zero matches, stating the queries and sources searched, so the coordinator can report 'no studies found' accurately.

Scenario 5: Claude Code for Continuous Integration (Questions 30-43)

You are integrating Claude Code into your Continuous Integration/Continuous Deployment (CI/CD) pipeline. The system runs automated code reviews, generates test cases, and provides feedback on pull requests. You need to design prompts that provide actionable feedback and minimize false positives.

Question 30

Select ONE answer.

Your CI job generates tests for new code. Currently one request produces tests that duplicate existing ones and miss untested branches. Which restructuring is best?

A. Generate tests only for functions that changed, ignoring what the test suite already contains.

B. A chain: analyze new code and existing tests for gaps, propose uncovered scenarios, generate tests for those, then run them and fix failures.

C. A single request telling Claude to generate as many tests as possible for every function it can find, so that coverage is as high as possible.

D. Generate tests twice with different seeds and merge both sets into the repository.

Question 31

Select ONE answer.

A PR changes the return type of normalizeAddress() in utils.ts, and a function in checkout.ts, untouched by the PR, relies on the old return type. Per-file review of the diff reports nothing. How should the integration pass be designed to catch this?

A. Add an instruction to each per-file pass to speculate about which other files might be affected.

B. Give the integration pass only the diff and ask it to think harder about possible side effects on the rest of the codebase and to list them in detail for the developer.

C. Rely on the unit tests to catch it and skip any cross-file analysis in the review.

D. Have the integration pass find callers of the changed interface (for example with Grep) and check those call sites against the new contract, even outside the diff.

Question 32

Select ONE answer.

A nightly job generates tests for a legacy repository. On the first night, the plan was fixed: tests for modules in alphabetical order. The agent discovers that module 'ledger' depends on a shared 'money' module with no tests, and ledger tests keep failing because of it. What is the best handling?

A. Adapt the plan: prioritize creating tests for the 'money' module first, then return to ledger, updating the plan as dependencies are discovered.

B. Restart the nightly job from scratch each time a dependency is found.

C. Continue in alphabetical order and accept the failing ledger tests until the 'money' module is eventually reached, since the plan was agreed in advance.

D. Skip any module that has failing tests and report them as untestable.

Question 33

Select ONE answer.

In CI, the MCP tool that posts review comments fails because the PR comes from a fork and the token has no write permission. It returns 'Failed to post'. The agent retries four times. What is the best tool design and agent behavior?

A. Return a transient error so the agent keeps retrying with backoff until the token is refreshed by an administrator or the job times out, whichever happens first.

B. Return a permission-category, non-retryable error with a clear description; the agent stops retrying and outputs the findings to the job summary instead.

C. Return success without posting, so that the CI check passes and nobody is blocked.

D. Return the stack trace of the HTTP client so that the model can work out the cause.

Question 34

Select ONE answer.

A CI job needs the Jira MCP server to read issue details with a token. How should the .mcp.json entry supply the token without committing it?

A. Add the token to the project's CLAUDE.md so the agent can read it when needed.

B. Ask the model to prompt a developer for the token each time the job runs.

C. Write the token as a literal string in .mcp.json, since the repository is private.

D. Reference an environment variable, such as ${JIRA_TOKEN}, and provide the value from the CI system's secret store.

Question 35

Select ONE answer.

A CI step runs 'claude "Review this PR"' and hangs, and when it works the output is free-form text that is hard to post as inline comments. The review should also follow the project's review criteria. Which combination fixes all three problems?

A. Run with -p and ask in the prompt for the output in a particular markdown layout, without any CLAUDE.md.

B. Run interactively with a here-document supplying answers to any prompts, and parse the text with regular expressions.

C. Run in non-interactive mode with -p, request --output-format json with a --json-schema for findings, and document review criteria in CLAUDE.md.

D. Use --output-format json without -p and put the review criteria in the user's ~/.claude/CLAUDE.md.

Question 36

Select ONE answer.

Before turning Claude's PR review into a blocking check, the team wants evidence about how good it is. What is the most informative evaluation?

A. Ask developers whether the comments seem helpful on their next ten PRs and decide by general impression, then switch on blocking if most say yes.

B. Run it on a labeled set of past PRs with known defects and known-clean changes, and measure precision and recall, including per category.

C. Count how many comments Claude posts per PR and treat higher counts as better coverage.

D. Run the review twice on each PR and measure only whether the two runs agree.

Question 37

Select ONE answer.

A team's pipeline runs tests, a Claude review, and then automatically deploys every merged change to production without human approval. Which change keeps automated build, test, and review on every PR while making production release a human decision?

A. Keep automated build, test, and review on every PR, but add a manual approval gate before production deployment.

B. Make the Claude review the approval gate that deploys to production when it finds no issues.

C. Keep everything as is, because deploying automatically after tests pass is what continuous integration means.

D. Remove automated testing from the pipeline, since CI is only about merging code frequently and tests slow down the integration of every change into the main branch.

Question 38

Select ONE answer.

The CI job that runs Claude Code only needs to read the repository and post review comments. It currently runs with a token that can push branches and change repository settings. What is the best change?

A. Use a token limited to reading code and posting PR comments, and keep branch and settings write access out of that job.

B. Keep the broad token and add a prompt instruction saying that the job must never push branches or edit repository settings under any circumstances.

C. Remove authentication from the job so that it fails if it tries to do anything sensitive.

D. Keep the broad token but log every action, and review the logs weekly.

Question 39

Select ONE answer.

To review generated code independently, a team passes the independent reviewer the diff plus the full chain of reasoning and commentary that the generating session produced. The reviewer mostly agrees with the generator. What should change?

A. Give the reviewer the diff, requirements, and project criteria, but not the generator's reasoning, so that it evaluates the change on its own.

B. Give the reviewer even more of the generator's reasoning and commentary, so that it understands the intent fully before judging the change and avoids false alarms.

C. Use the same session as the generator for the review, with a prompt telling it to be strict.

D. Ask the reviewer to disagree with the generator at least twice per review.

Question 40

Select ONE answer.

To reduce noisy findings, a team runs the full review three times and posts only findings that appear in at least two runs. After a month, real bugs found in only one run are being missed. What is the better approach to precision?

A. Keep the voting and lower the severity of all findings to hide the missed ones.

B. Give the review explicit criteria for what to report and skip and use focused per-file passes, instead of requiring agreement among runs.

C. Require agreement in all three runs instead of two, and drop single-run findings entirely, for even stronger filtering of the noisy findings that developers complain about.

D. Run the review five times and post the union of all findings.

Question 41

Select ONE answer.

Developers distrust the review because it flags minor style preferences and local naming habits as problems, mixed in with real bugs. The prompt says 'only report high-confidence findings'. Which change best improves precision?

A. State the categories to report (bugs, security issues, correctness risks) and the categories to skip (minor style, locally accepted patterns), not confidence.

B. Add the instruction 'be conservative' to the end of the prompt.

C. Strengthen the confidence instruction to 'only report findings you are 95% sure about, and say nothing about anything less certain than that in the diff'.

D. Ask the model to rank all findings and post only the top three per PR.

Question 42

Select ONE answer.

A team wants a review pipeline for large PRs that balances depth, cross-file coverage, and trust in blocking findings. Which design is most sound?

A. Per-file passes only, with all findings posted directly as blocking comments.

B. One pass over the whole PR, followed by a second pass by the same session to double-check and refine all of its own findings before posting.

C. Per-file local passes, then a cross-file integration pass, then an independent verification pass on any finding that would block the merge.

D. Three parallel full-PR passes with majority voting and no integration pass.

Question 43

Select ONE answer.

A team reports that its Claude review is 'right 92% of the time' and wants to make all categories blocking. Why is this insufficient evidence?

A. An overall figure can hide poor accuracy in particular categories, so precision should be measured per finding category before choosing which, if any, to block on.

B. The figure is sufficient if the review sample is large, because a big sample makes per-category breakdowns unnecessary.

C. Blocking depends only on how fast the review runs, not on its accuracy.

D. The figure is insufficient only if it was computed by developers instead of by the model.

Scenario 6: Structured Data Extraction (Questions 44-60)

You are building a structured data extraction system using Claude. The system extracts information from unstructured documents, validates the output using JSON schemas, and maintains high accuracy. It must handle edge cases gracefully and integrate with downstream systems.

Question 44

Select ONE answer.

Compliance wants every write the extraction agent makes to the ledger to be traceable to the exact source document and extraction run, regardless of whether the model remembers to record this. Which implementation is most reliable?

A. A prompt instruction asking the agent to include a source document ID and a run ID in the free-text note of every write it makes to the ledger.

B. A pre-call hook on the write tool that records the record hash, source document ID, and run ID to an audit log before allowing the call.

C. A weekly script that reconstructs which document produced each ledger entry by matching amounts.

D. An optional field in the schema where the model may add a reference to the source if it chooses.

Question 45

Select ONE answer.

A document-management MCP tool returns status as numeric codes (1 received, 2 processed, 3 rejected). The extraction agent sometimes treats 3 as a success because it appears larger. What is the best fix?

A. A PostToolUse hook that maps numeric codes to labeled statuses before the model sees the result.

B. A request to the tool's owners to renumber the codes so that failures are small numbers.

C. A system prompt paragraph that says higher status codes are worse and should be treated as failures.

D. An example showing a status code 3 being handled correctly, placed in the extraction prompt.

Question 46

Select ONE answer.

In a batch of 500 documents, a bug in a new prompt version causes 40% of extractions to fail validation. The pipeline continues posting the 60% that passed to the ledger, and the review team is overwhelmed. What deterministic safeguard fits best?

A. A circuit-breaker hook that halts automated posting and redirects the batch to review when the validation failure rate crosses a set threshold.

B. A rule that retries each failed document indefinitely until it validates.

C. A nightly cleanup job that removes bad ledger entries after the batch finishes.

D. A prompt instruction telling the agent to slow down and become more careful when it notices that many of the documents it processes fail validation.

Question 47

Select ONE answer.

The extraction agent occasionally needs a currency exchange rate. The finance MCP server exposes 17 tools, including payment and ledger operations. What is the best way to provide this capability?

A. Connect the agent to all 17 finance tools so that it can handle any financial need it encounters without asking for new access later.

B. Ask the agent to estimate exchange rates from its general knowledge, which is usually close.

C. Give the agent a narrow get_fx_rate tool for the one lookup it needs, and keep the rest of the finance tools out of its tool set.

D. Pass the rate table in every prompt regardless of whether a conversion is needed.

Question 48

Select ONE answer.

The extraction agent currently holds both extract_invoice and post_to_ledger. An incorrect extraction was posted before validation. The team wants to prevent the agent from posting at all. What is the best structure?

A. Keep post_to_ledger but set tool_choice to 'auto' so the agent uses it only when needed.

B. Keep post_to_ledger and rename it submit_for_review so that the agent treats it more carefully.

C. Remove post_to_ledger from the agent's tools, and have deterministic pipeline code post only records that have passed validation.

D. Keep post_to_ledger and instruct the agent to call it only after the user confirms the extraction, as the final step of each document.

Question 49

Select ONE answer.

The ledger API rejects a posting with HTTP 422 and a list of field errors, for example 'cost_center: unknown code'. The MCP wrapper converts this into 'Posting failed'. The agent retries unchanged. What is the best wrapper behavior?

A. Swallow the error and report success to avoid blocking the pipeline.

B. Return 'Posting failed' with a hint that the agent should try again after a short wait.

C. Return a transient error so that the retry logic handles it automatically, since most 4xx errors clear after a while and the posting will then go through.

D. Return an isError result with category validation, isRetryable false, and the field-level errors, so the agent or pipeline can correct the input.

Question 50

Select ONE answer.

extract_invoice returns amounts as integers such as 125000, but the description says only 'amount of the invoice'. Downstream agents sometimes treat the value as dollars and sometimes as cents. How should the tool be improved?

A. Return the amount as a formatted string such as '$1,250.00' for readability.

B. Leave the contract unchanged and ask downstream agents to infer the unit from magnitude.

C. State in the description and schema that amounts are integers in minor currency units (cents) alongside a separate ISO currency code field.

D. Add a sentence to the description saying that the amount is 'accurate', 'verified', and 'should be used carefully' by every downstream agent that reads it.

Question 51

Select ONE answer.

A validation-retry loop sends specific errors back to the model for up to five rounds. For some documents it still fails every time, and the retries consume budget. What is the best bounded design?

A. Raise the limit to twenty rounds so that the model has more chances to get it right, with the validation error attached on every round, whatever the cost.

B. Limit retries to a small number (for example two), then route the document to human review with the document, the attempts, and the errors attached.

C. Drop documents that fail validation twice, to keep the pipeline moving.

D. After the final failure, fill the failing fields with defaults and accept the record.

Question 52

Select ONE answer.

A team wants to use the Message Batches API for 20,000 documents, but their current extraction flow calls a vendor lookup tool in the middle of each request and continues with the result. What is the correct adaptation?

A. Switch the whole job to the synchronous API since any tool use rules out batch processing.

B. Embed the entire vendor database in each request to avoid needing a lookup.

C. Submit the batch as is, and expect each request to call the lookup tool, receive the vendor data, and continue its extraction in the very same request, as it does today in the synchronous flow.

D. Restructure to a single-turn extraction per batch request, and do vendor lookups in pipeline code before or after the batch, since batch requests can't run tools mid-request.

Question 53

Select ONE answer.

Prices in documents like '1,299.99' are extracted into a floating-point field. Totals computed later sometimes differ by a cent from the invoice's printed total. What schema design best protects monetary precision?

A. Store amounts as free text and let each downstream consumer parse them as it prefers.

B. Keep floats but tell the model in the prompt instructions to round every value to two decimals before it fills in the field, so the cents are right.

C. Represent amounts as integers in minor currency units, or as decimal strings, with a separate currency code, and do arithmetic in decimal types.

D. Use floats but multiply by 1.0001 to offset the observed difference.

Question 54

Select ONE answer.

A batch of 100 documents returns 12 failures: 3 exceeded context limits, and 9 returned output that failed schema validation. What is the best next step?

A. Switch all 100 documents to the synchronous API and retry them one at a time.

B. Using custom_id, chunk the 3 oversized documents and resubmit them, and refine the prompt for the 9 validation failures, resubmitting only those.

C. Resubmit all 100 documents unchanged in a new batch and hope that the failures were transient, since the previous run may have hit temporary limits.

D. Resubmit only the 3 oversized documents and discard the 9 validation failures as noise.

Question 55

Select ONE answer.

Which extraction design best detects inconsistent numbers and inconsistent source data and sends the right records to review?

A. Extract only the stated total and trust it, because documents are generally correct and a second total only adds confusion for reviewers downstream.

B. Ask the model to say 'I am confident' or 'I am not confident' at the end of its response.

C. Extract only the calculated total from line items and discard the stated total.

D. Extract stated_total and calculated_total, add a conflict_detected boolean when the source itself disagrees, and route flagged records to review.

Question 56

Select ONE answer.

A customer profile is built by merging fields from three documents (contract, invoice, onboarding form). Reviewers can't tell which document supplied each field, or when two documents disagreed. What is the best design?

A. Store each field with its source document, location, excerpt, and date, and annotate conflicts with the competing values and sources instead of silently picking one.

B. Keep the most recently dated value for every field and discard the others, so that the profile stays clean and reviewers see only one value per field, which is also the simplest design.

C. Have the model write a paragraph summarizing the profile, with sources named in passing.

D. Add a single note listing all three documents as sources for the whole profile.

Question 57

Select exactly TWO answers (choose 2).

An analyst report built from extracted data includes quarterly financial figures, a set of news developments, and a list of technical findings about system performance. Select the TWO formatting choices that best suit the content types.

A. Present the quarterly financial figures as a table.

B. Present the news developments as a table with one cell per sentence.

C. Present the technical findings as a structured list with their sources.

D. Render all content as one uniform bulleted list for consistency.

E. Convert the financial figures into a flowing narrative paragraph to make the report more engaging.

Question 58

Select ONE answer.

Why sample high-confidence extractions at random for human review when you already review all low-confidence ones?

A. To measure the real error rate among records that the system trusts and to detect new error patterns that confidence scores do not reveal.

B. To reduce the number of low-confidence records that need review by retraining the reviewers, which management prefers.

C. To confirm that the confidence threshold can be raised without checking any real outcomes.

D. To avoid having to calibrate confidence thresholds on a labeled set.

Question 59

Select ONE answer.

In a nightly run, 8 of 100 extracted records are rejected by the downstream system for various reasons. The pipeline currently rolls back all 100 and reports a generic failure. What is the better behavior?

A. Commit the 92 valid records, report the 8 failures with their reasons for correction or review, and keep the rest of the run's results.

B. Commit the 92 and silently drop the 8 so the run appears to succeed.

C. Keep rolling back all 100 so the downstream system is always consistent with a single run.

D. Retry the whole run of 100 repeatedly until all pass.

Question 60

Select ONE answer.

An extraction returns total = $4,200 with field confidence 0.97, but the document contains two different totals ($4,200 and $4,020) and conflict_detected is true. The routing rule auto-accepts anything above 0.95 confidence. What should the routing logic do?

A. Auto-accept it, since 0.97 exceeds the threshold and thresholds are the agreed rule.

B. Route it to human review because contradictory source data warrants review regardless of the confidence score.

C. Reject it permanently because documents with conflicting totals are invalid.

D. Auto-accept it but lower the confidence to 0.90 so that dashboards show uncertainty.

End of examination. Stop here, then check your answers using the key that follows.

5. Quick answer key and score tracker

Mark each question you answered correctly (all required selections, no wrong ones). Then complete the domain tracker below.

No.	Ans.	Domain	Scen.	Mine
1	**D**	D1	S2	
2	**A**	D1	S2	
3	**C**	D3	S2	
4	**B**	D3	S2	
5	**B**	D3	S2	
6	**A,C**	D3	S2	
7	**A,E**	D3	S2	
8	**A,E**	D3	S2	
9	**A**	D3	S2	
10	**C**	D3	S2	
11	**B**	D4	S2	
12	**C**	D4	S2	
13	**A**	D5	S2	
14	**B**	D1	S3	
15	**C**	D1	S3	
16	**B**	D1	S3	
17	**D**	D1	S3	
18	**D**	D1	S3	
19	**C**	D1	S3	
20	**D**	D1	S3	
21	**D**	D1	S3	
22	**D**	D2	S3	
23	**C**	D2	S3	
24	**D**	D2	S3	
25	**A,D**	D2	S3	
26	**B**	D2	S3	
27	**A,B,F**	D4	S3	
28	**C**	D5	S3	
29	**D**	D5	S3	
30	**B**	D1	S5	

No.	Ans.	Domain	Scen.	Mine
31	**D**	D1	S5	
32	**A**	D1	S5	
33	**B**	D2	S5	
34	**D**	D2	S5	
35	**C**	D3	S5	
36	**B**	D3	S5	
37	**A**	D3	S5	
38	**A**	D3	S5	
39	**A**	D4	S5	
40	**B**	D4	S5	
41	**A**	D4	S5	
42	**C**	D4	S5	
43	**A**	D5	S5	
44	**B**	D1	S6	
45	**A**	D1	S6	
46	**A**	D1	S6	
47	**C**	D2	S6	
48	**C**	D2	S6	
49	**D**	D2	S6	
50	**C**	D2	S6	
51	**B**	D4	S6	
52	**D**	D4	S6	
53	**C**	D4	S6	
54	**B**	D4	S6	
55	**D**	D4	S6	
56	**A**	D5	S6	
57	**A,C**	D5	S6	
58	**A**	D5	S6	
59	**A**	D5	S6	
60	**B**	D5	S6	

Domain-level score tracker

Domain	Questions	Number list	My correct	% correct
1. Agentic Architecture & Orchestration	16	1, 2, 14, 15, 16, 17, 18, 19, 20, 21, 30, 31, 32, 44, 45, 46		
2. Tool Design & MCP Integration	11	22, 23, 24, 25, 26, 33, 34, 47, 48, 49, 50		
3. Claude Code Configuration & Workflows	12	3, 4, 5, 6, 7, 8, 9, 10, 35, 36, 37, 38		
4. Prompt Engineering & Structured Output	12	11, 12, 27, 39, 40, 41, 42, 51, 52, 53, 54, 55		
5. Context Management & Reliability	9	13, 28, 29, 43, 56, 57, 58, 59, 60		
Total	**60**			

Interpretation: percent-correct here is a practice signal only. The real exam's pass/fail decision is based on a single scaled score (cut score 720 of 1,000); domain percentages on the real score report are informational.

6. Detailed answer key and explanations

Each entry gives the correct answer, why it is correct, why every other option is wrong, the concept tested, the detail in the scenario that decides the answer, and a production lesson.

Q1 | Domain 1: Agentic Architecture & Orchestration | Scenario 2: Code Generation with Claude Code

Task statement 1.7 | Intermediate | Single answer

Correct answer: D - Confirmed conclusions, decisions already made, files of interest, and open questions, leaving out raw tool output tied to the old runner.

Why the correct answer is correct: When prior tool results are stale, a new session seeded with a structured summary is more reliable than resuming. The summary should carry durable facts and decisions and leave out outdated raw output.

Option-by-option:

A. **[INCORRECT]** Only the first and last message of the old session, since they contain the goal and the final result.
Middle messages contain the findings and decisions that should survive.

B. **[INCORRECT]** Every command the old session ran, in order, so that the new session can replay them and compare results.
Replaying commands is slow and repeats exploration the summary is meant to avoid.

C. **[INCORRECT]** The complete transcript of the old session, pasted in full, so that nothing the previous session saw is lost and the new session has every detail available.
A full paste brings the stale tool results back into the new context.

D. **[CORRECT]** Confirmed conclusions, decisions already made, files of interest, and open questions, leaving out raw tool output tied to the old runner.
Correct. The summary carries what is still valid and omits stale output.

Key concept: Fresh session with injected summary (Task 1.7).
Exam reasoning: The runner upgrade invalidated earlier outputs but not the conclusions about test design.
Production lesson: Carry forward conclusions, not stale evidence.

Q2 | Domain 1: Agentic Architecture & Orchestration | Scenario 2: Code Generation with Claude Code

Task statement 1.7 | Foundational | Single answer

Correct answer: A - Comparing two candidate refactoring approaches, starting from the same completed analysis of a module.

Why the correct answer is correct: fork_session creates independent branches from a shared analysis baseline to explore divergent approaches.

Option-by-option:

A. **[CORRECT]** Comparing two candidate refactoring approaches, starting from the same completed analysis of a module.
Correct. Forks explore divergent approaches from a shared baseline.

B. **[INCORRECT]** Recovering from a crash by restoring the last saved state of every tool call.
Crash recovery relies on persisted state, not branching.

C. **[INCORRECT]** Continuing yesterday's debugging work on the same bug with the same accumulated context.
Continuing a prior conversation is session resumption.

D. **[INCORRECT]** Beginning an unrelated task in a different repository with no connection to the earlier analysis.
An unrelated task has no shared baseline to fork from.

Key concept: fork_session use cases (Task 1.7).
Exam reasoning: Only one option has a shared baseline and divergent approaches.
Production lesson: Fork to compare, resume to continue.

Q3 | Domain 3: Claude Code Configuration & Workflows | Scenario 2: Code Generation with Claude Code

Task statement 3.1 | Intermediate | Single answer

Correct answer: C - (a) user-level ~/.claude/CLAUDE.md, (b) directory-level CLAUDE.md in payments/, (c) project-level CLAUDE.md, (d) a .claude/rules/ file scoped by a **/*.tf glob.

Why the correct answer is correct: CLAUDE.md supports user, project, and directory scopes; path-scoped rules in .claude/rules/ apply conventions by file pattern, regardless of directory.

Option-by-option:

A. **[INCORRECT]** (a) a skill, (b) the project-level CLAUDE.md, (c) the user-level ~/.claude/CLAUDE.md, (d) a directory-level CLAUDE.md in terraform/.
Repository-wide naming standards in a personal file are not shared, and one directory file misses .tf files elsewhere.

B. **[INCORRECT]** (a) directory-level CLAUDE.md in payments/, (b) the project-level CLAUDE.md, (c) a .claude/rules/ file without paths, (d) the user-level ~/.claude/CLAUDE.md.
This scrambles scope for all four items.

C. **[CORRECT]** (a) user-level ~/.claude/CLAUDE.md, (b) directory-level CLAUDE.md in payments/, (c) project-level CLAUDE.md, (d) a .claude/rules/ file scoped by a **/*.tf glob.
Correct. Each item uses the narrowest mechanism that reaches the right audience and files.

D. **[INCORRECT]** (a) project-level CLAUDE.md, (b) user-level ~/.claude/CLAUDE.md, (c) directory-level CLAUDE.md in each package, (d) a skill invoked manually before editing any Terraform file.
This puts personal preferences in the shared file and relies on manual invocation for Terraform.

Key concept: Configuration scope mapping (Task 3.1, 3.3).
Exam reasoning: Each item differs in audience and file reach, which determines its mechanism.
Production lesson: Choose scope by who and which files.

Q4 | Domain 3: Claude Code Configuration & Workflows | Scenario 2: Code Generation with Claude Code

Task statement 3.2 | Foundational | Single answer

Correct answer: B - Skills, defined as SKILL.md files under .claude/skills/.

Why the correct answer is correct: Skills in .claude/skills/ use SKILL.md frontmatter supporting context: fork, allowed-tools, and argument-hint.

Option-by-option:

A. **[INCORRECT]** The settings file for hooks, as fields on the PostToolUse entry.
These are not hook settings.

B. **[CORRECT]** Skills, defined as SKILL.md files under .claude/skills/.
Correct. These frontmatter options belong to skills.

C. **[INCORRECT]** The project-level CLAUDE.md file, in a block at the top of the file.
CLAUDE.md holds instructions, not these skill options.

D. **[INCORRECT]** The .mcp.json file, as properties on each server entry.
.mcp.json configures MCP servers.

Key concept: Skill frontmatter (Task 3.2).
Exam reasoning: Only skills use these options.
Production lesson: Know where each option lives.

Q5 | Domain 3: Claude Code Configuration & Workflows | Scenario 2: Code Generation with Claude Code

Task statement 3.5 | Intermediate | Single answer

Correct answer: B - (i) provide two or three concrete input/output examples, (ii) use the interview pattern, (iii) write tests first and iterate on failures.

Why the correct answer is correct: Concrete examples fix inconsistent interpretation; the interview pattern surfaces unconsidered considerations in unfamiliar domains; test-driven iteration gives verifiable feedback on edge cases.

Option-by-option:

A. **[INCORRECT]** (i) write tests first and iterate on failures, (ii) provide concrete input/output examples, (iii) use the interview pattern.
Examples need domain knowledge to write, and an interview doesn't define edge case behavior.

B. **[CORRECT]** (i) provide two or three concrete input/output examples, (ii) use the interview pattern, (iii) write tests first and iterate on failures.
Correct. Each technique matches the problem it was designed for.

C. **[INCORRECT]** (i) use the interview pattern, (ii) write tests first and iterate on failures, (iii) provide two or three concrete input/output examples for each edge case.
Tests can't be written well in an unfamiliar domain, and the interview doesn't fix inconsistent interpretation.

D. **[INCORRECT]** (i) repeat the prose with greater emphasis, (ii) provide one long example, (iii) ask Claude to rate its confidence.
None of these addresses the underlying problems.

Key concept: Iterative refinement techniques (Task 3.5).

Exam reasoning: Each problem has a technique named for it.
Production lesson: Match the technique to the failure.

Q6 | Domain 3: Claude Code Configuration & Workflows | Scenario 2: Code Generation with Claude Code

Task statement 3.3 | Intermediate | Select 2

Correct answer: A, C - Path-scoped rules load only when Claude works on matching files, which reduces irrelevant context and token usage. / A rule with paths: ["**/*.test.*"] applies to test files wherever they live in the repository.

Why the correct answer is correct: YAML frontmatter paths fields contain glob patterns for conditional rule activation, loading only for matching files.

Option-by-option:

A. **[CORRECT]** Path-scoped rules load only when Claude works on matching files, which reduces irrelevant context and token usage.
Correct. Conditional loading is the main advantage.

B. **[INCORRECT]** A paths value can only be an exact file name, so a glob pattern isn't allowed.
Glob patterns are supported.

C. **[CORRECT]** A rule with paths: ["**/*.test.*"] applies to test files wherever they live in the repository.
Correct. A recursive glob matches by file pattern, not by directory.

D. **[INCORRECT]** A directory-level CLAUDE.md automatically applies to matching files in sibling directories.
Directory-level files are bound to their own directories.

E. **[INCORRECT]** Path-scoped rules only apply after a developer runs a slash command naming the rule file.
Rules apply automatically based on matching paths.

Key concept: Path-scoped rules (Task 3.3).
Exam reasoning: Two statements describe actual behavior; the others invent constraints.
Production lesson: Rules follow file patterns.

Q7 | Domain 3: Claude Code Configuration & Workflows | Scenario 2: Code Generation with Claude Code

Task statement 3.4 | Foundational | Select 2

Correct answer: A, E - Adding a date validation conditional to one existing function with fully specified behavior. / Fixing a crash where the stack trace points at one specific line in one function.

Why the correct answer is correct: Direct execution suits simple, well-scoped changes such as single-function bug fixes; plan mode suits large-scale, multi-file, or architectural work.

Option-by-option:

A. **[CORRECT]** Adding a date validation conditional to one existing function with fully specified behavior.
Correct. It is a well-scoped single-function change.

B. **[INCORRECT]** Choosing between two integration approaches that need different infrastructure.
Multiple valid approaches with different consequences call for plan mode.

C. **[INCORRECT]** Restructuring a monolith into microservices, which requires decisions about service boundaries.
This is a large architectural task.

D. **[INCORRECT]** Migrating a logging library across roughly 45 files with several custom wrappers.
A broad migration benefits from planning.

E. **[CORRECT]** Fixing a crash where the stack trace points at one specific line in one function.
Correct. Scope is clear and local.

Key concept: Mode selection (Task 3.4).
Exam reasoning: The two correct choices are local and well specified; the other three are broad or involve choices.
Production lesson: Small and clear: just do it.

Q8 | Domain 3: Claude Code Configuration & Workflows | Scenario 2: Code Generation with Claude Code

Task statement 3.2 | Intermediate | Select 2

Correct answer: A, E - allowed-tools in the frontmatter restricts which tools the skill may use while it runs. / context: fork runs a skill in an isolated sub-agent context so verbose output doesn't pollute the main conversation.

Why the correct answer is correct: Skill frontmatter provides context: fork for isolation, allowed-tools for restriction, and argument-hint for parameter prompts; user-scoped files are personal.

Option-by-option:

A. **[CORRECT]** allowed-tools in the frontmatter restricts which tools the skill may use while it runs.
Correct. It can limit a skill to safe operations.

B. **[INCORRECT]** Commands in ~/.claude/commands/ are shared with teammates through version control.
User-scoped commands are personal.

C. **[INCORRECT]** Skills in ~/.claude/skills/ replace the project's skill of the same name for the whole team.
Personal skills affect only the individual, and distinct names avoid collisions.

D. **[INCORRECT]** argument-hint restricts which tools a skill is allowed to use.
argument-hint prompts for parameters; it doesn't restrict tools.

E. **[CORRECT]** context: fork runs a skill in an isolated sub-agent context so verbose output doesn't pollute the main conversation.
Correct. That is the purpose of context: fork.

Key concept: Skills and commands (Task 3.2).
Exam reasoning: Two statements correctly assign purposes to options; the others swap them or misstate scope.
Production lesson: Learn what each frontmatter field actually does.

Q9 | Domain 3: Claude Code Configuration & Workflows | Scenario 2: Code Generation with Claude Code

Task statement 3.3 | Foundational | Single answer

Correct answer: A - When it is working with files that match terraform/**/*, such as an edit to terraform/network/main.tf.

Why the correct answer is correct: Path-scoped rules load only when working on files matching the glob patterns in the paths field.

Option-by-option:

A. **[CORRECT]** When it is working with files that match terraform/**/*, such as an edit to terraform/network/main.tf.
Correct. Path-scoped rules load for matching files.

B. **[INCORRECT]** At the start of every session, regardless of which files are being edited.
That is how unscoped instructions behave.

C. **[INCORRECT]** Only for files directly inside terraform/, not for files in subdirectories.
*The ** pattern includes subdirectories.*

D. **[INCORRECT]** Only when the developer types the rule's file name as a slash command.
Rules are not invoked manually.

Key concept: Conditional rule loading (Task 3.3).
Exam reasoning: The glob terraform/**/* matches files at any depth under terraform/.
Production lesson: Scope decides when it loads.

Q10 | Domain 3: Claude Code Configuration & Workflows | Scenario 2: Code Generation with Claude Code

Task statement 3.5 | Intermediate | Single answer

Correct answer: C - Write tests for correctness, edge cases, and the 200 ms requirement first, then have Claude implement and iterate by sharing failing test output.

Why the correct answer is correct: Test-driven iteration writes test suites covering expected behavior, edge cases, and performance requirements first, then iterates by sharing test failures.

Option-by-option:

A. **[INCORRECT]** Describe the performance requirement in more detail and ask Claude to reason carefully about algorithmic complexity before writing any code, then ship it.
Reasoning about complexity alone doesn't verify actual performance.

B. **[INCORRECT]** Implement first, then ask Claude to produce tests that pass for the implementation.
Tests derived from the implementation won't catch missing requirements.

C. **[CORRECT]** Write tests for correctness, edge cases, and the 200 ms requirement first, then have Claude implement and iterate by sharing failing test output.
Correct. Tests make correctness and performance requirements verifiable.

D. **[INCORRECT]** Ask for three different implementations and choose the one that looks the most efficient.
Appearance isn't measurement.

Key concept: Test-driven iteration (Task 3.5).
Exam reasoning: Performance and edge cases can be checked by tests.
Production lesson: Make the requirements executable.

Q11 | Domain 4: Prompt Engineering & Structured Output | Scenario 2: Code Generation with Claude Code

Task statement 4.2 | Intermediate | Single answer

Correct answer: B - A few examples of ambiguous cases showing which test type was chosen and why it was preferred over the alternative.

Why the correct answer is correct: Few-shot examples for ambiguous scenarios that show reasoning for why one action was chosen over plausible alternatives improve consistency and judgment.

Option-by-option:

A. **[INCORRECT]** A reminder to think carefully before choosing and to be consistent from file to file.
A reminder provides no criteria.

B. **[CORRECT]** A few examples of ambiguous cases showing which test type was chosen and why it was preferred over the alternative.
Correct. Reasoned examples teach the decision boundary.

C. **[INCORRECT]** A rule that every behavior gets both a unit test and an integration test to avoid choosing.
Doubling tests wastes effort and produces redundancy.

D. **[INCORRECT]** An instruction to pick whichever test type is faster to write in each case.
Speed isn't the appropriate criterion.

Key concept: Few-shot for ambiguous decisions (Task 4.2).
Exam reasoning: The ambiguity is the point, and reasoned examples address it.
Production lesson: Show how to decide, not just what to do.

Q12 | Domain 4: Prompt Engineering & Structured Output | Scenario 2: Code Generation with Claude Code

Task statement 4.2 | Challenging | Single answer

Correct answer: C - Replace the examples with ones drawn from current, preferred code written in async/await, since examples strongly shape generated style.

Why the correct answer is correct: Few-shot examples shape output closely, so they must demonstrate acceptable, current patterns; examples of outdated patterns propagate them.

Option-by-option:

A. **[INCORRECT]** Add more callback examples to improve consistency, then rely on a code formatter to rewrite them.
A formatter cannot convert callback structures to async/await.

B. **[INCORRECT]** Keep the callback examples and add a sentence to the prompt saying that async/await is the preferred style for all new code written from now on.
Contradictory examples and instructions give unreliable results; replacing the examples addresses the cause.

C. **[CORRECT]** Replace the examples with ones drawn from current, preferred code written in async/await, since examples strongly shape generated style.
Correct. Examples should reflect the desired patterns.

D. **[INCORRECT]** Remove all examples so that the model uses its own default style instead.
The default style may not match the house style either.

Key concept: Example quality (Task 4.2).
Exam reasoning: Generated code mirrors the examples' style, indicating that the examples are the cause.
Production lesson: Examples are instructions. Curate them.

Q13 | Domain 5: Context Management & Reliability | Scenario 2: Code Generation with Claude Code

Task statement 5.4 | Foundational | Single answer

Correct answer: A - It reduces context usage by condensing the conversation so a long session can continue.

Why the correct answer is correct: Use /compact to reduce context usage during extended exploration sessions when context fills with verbose discovery output.

Option-by-option:

A. **[CORRECT]** It reduces context usage by condensing the conversation so a long session can continue.
Correct. /compact is used when context fills with verbose discovery output.

B. **[INCORRECT]** It permanently deletes the project's CLAUDE.md to save space in the repository.
/compact doesn't delete configuration files.

C. **[INCORRECT]** It compresses source files on disk so that they take less space.
It acts on conversation context, not on the file system.

D. **[INCORRECT]** It restarts Claude Code with a different model that has a larger context window.
It doesn't change the model.

Key concept: /compact (Task 5.4).
Exam reasoning: The command targets conversation context.
Production lesson: Compact context, keep durable notes elsewhere.

Q14 | Domain 1: Agentic Architecture & Orchestration | Scenario 3: Multi-Agent Research System

Task statement 1.1 | Foundational | Single answer

Correct answer: B - Execute the requested tool(s), append the tool results to the conversation, and send another request.

Why the correct answer is correct: The agentic loop continues when stop_reason is 'tool_use' by executing tools and returning results, and terminates when it is 'end_turn'.

Option-by-option:

A. **[INCORRECT]** Discard the tool request and call the model again with the same messages to get a different response.
That ignores the request and wastes a call.

B. **[CORRECT]** Execute the requested tool(s), append the tool results to the conversation, and send another request.
Correct. That continues the loop with new information.

C. **[INCORRECT]** Return the response text to the coordinator as the final answer, since the model has finished.
'tool_use' means the model is waiting for tool results, not that it finished.

D. **[INCORRECT]** Stop the loop and log an error because the subagent should not use tools.
Tool use is the normal path for an agent that has tools.

Key concept: Loop control flow (Task 1.1).
Exam reasoning: 'tool_use' is the signal to run tools and continue.
Production lesson: Know the two key stop reasons cold.

Q15 | Domain 1: Agentic Architecture & Orchestration | Scenario 3: Multi-Agent Research System

Task statement 1.1 | Intermediate | Single answer, secondary: D5

Correct answer: C - Keep stop_reason as the normal terminator, and add token and tool-call budgets that stop gracefully and produce a partial report stating what was not covered.

Why the correct answer is correct: stop_reason should drive normal termination. Budget limits are safety backstops, and when triggered they should degrade gracefully with disclosed gaps rather than silently truncate.

Option-by-option:

A. **[INCORRECT]** Use a hard iteration cap of 5 as the primary stopping rule, and return whatever has been gathered by then as the final report for the user without comment.
A small cap as the primary mechanism cuts legitimate research short without disclosure.

B. **[INCORRECT]** Stop the loop whenever the model writes a sentence beginning 'In summary', since that signals completion.
Parsing phrases is an anti-pattern.

C. **[CORRECT]** Keep stop_reason as the normal terminator, and add token and tool-call budgets that stop gracefully and produce a partial report stating what was not covered.
Correct. Budgets protect cost, and graceful degradation avoids silent truncation.

D. **[INCORRECT]** Remove all limits and rely on the model to decide when it has searched enough.
Without limits, rare runaway runs can cost a great deal.

Key concept: Loop termination with backstops (Task 1.1, 5.3).
Exam reasoning: Both cost control and honesty about coverage are required.
Production lesson: Guards should fail visibly.

Q16 | Domain 1: Agentic Architecture & Orchestration | Scenario 3: Multi-Agent Research System

Task statement 1.3 | Intermediate | Single answer, secondary: D5

Correct answer: B - The research question and quality criteria, the findings as structured records (claim, excerpt, source, date), noted conflicts, and the required format.

Why the correct answer is correct: Subagent prompts must contain the needed context explicitly, ideally as structured data separating content from metadata (source, date) plus goals and quality criteria.

Option-by-option:

A. **[INCORRECT]** A short summary of the findings in prose, without sources, to keep the prompt compact.
Dropping sources loses attribution for the final report.

B. **[CORRECT]** The research question and quality criteria, the findings as structured records (claim, excerpt, source, date), noted conflicts, and the required format.
Correct. It provides goals, evidence with provenance, and expectations.

C. **[INCORRECT]** The research question and the instruction 'summarize everything found so far', with no findings included in the prompt.
Subagents do not inherit the coordinator's context, so without the findings in the prompt it has nothing to summarize.

D. **[INCORRECT]** The complete raw text of all fetched pages, with no structure, so that no detail is lost.
Unstructured bulk is hard to use and consumes context.

Key concept: Context packets (Task 1.3, 5.6).
Exam reasoning: Only one option equips the subagent to write a cited, conflict-aware synthesis.
Production lesson: A subagent can only use what it is handed.

Q17 | Domain 1: Agentic Architecture & Orchestration | Scenario 3: Multi-Agent Research System

Task statement 1.2 | Intermediate | Single answer

Correct answer: D - Partition the scope: assign distinct subtopics or source types to each subagent, such as regions, regulatory bodies, and industry commentary.

Why the correct answer is correct: Partition research scope across subagents to minimize duplication, for example by assigning distinct subtopics or source types.

Option-by-option:

A. **[INCORRECT]** Keep the same assignment but have the synthesis subagent remove duplicate results later.
Deduplication afterward doesn't fill the coverage gap.

B. **[INCORRECT]** Replace the two subagents with a single one that researches the whole topic sequentially.
This loses parallelism and risks narrow coverage.

C. **[INCORRECT]** Ask both subagents to work harder and to be more thorough, in the hope that they will find different material by chance on the next run, without changing assignments.
Without distinct assignments, overlap persists.

D. **[CORRECT]** Partition the scope: assign distinct subtopics or source types to each subagent, such as regions, regulatory bodies, and industry commentary.
Correct. Partitioning minimizes duplication and widens coverage.

Key concept: Scope partitioning (Task 1.2).
Exam reasoning: Overlap with thin coverage elsewhere indicates unpartitioned assignment.
Production lesson: Give each worker its own lane.

Q18 | Domain 1: Agentic Architecture & Orchestration | Scenario 3: Multi-Agent Research System

Task statement 1.1 | Challenging | Single answer

Correct answer: D - Have the hook return a tool_result explaining that the call was blocked and why, so the model can continue with another source or note the gap.

Why the correct answer is correct: Hooks enforce rules deterministically, and their outcomes should flow back into the loop as tool results so that the model can recover or note the limitation.

Option-by-option:

A. **[INCORRECT]** Remove the hook, since blocking domains should be handled by telling the model in the prompt.
The prompt cannot guarantee compliance.

B. **[INCORRECT]** Silently return an empty result for the blocked call so the model assumes there is nothing at that domain.
Silent empty results mislead the model about the cause.

C. **[INCORRECT]** Let the exception end the loop, because a blocked call should always abort the whole task for safety reasons and report the policy violation to the user.
Aborting loses all other work.

D. **[CORRECT]** Have the hook return a tool_result explaining that the call was blocked and why, so the model can continue with another source or note the gap.
Correct. The loop continues, and the model can adapt.

Key concept: Hook outcomes in the agentic loop (Task 1.1, 1.5).
Exam reasoning: The loop ended unnecessarily, and the model never learned why.
Production lesson: Make blocks informative, not fatal.

Q19 | Domain 1: Agentic Architecture & Orchestration | Scenario 3: Multi-Agent Research System

Task statement 1.2 | Intermediate | Single answer

Correct answer: C - Run web search and document analysis in parallel, then run synthesis once both complete, then report writing.

Why the correct answer is correct: Independent subagent tasks can run in parallel through multiple Task calls in one response; dependent steps run after their inputs are available.

Option-by-option:

A. **[INCORRECT]** Run all four steps in parallel and merge the results at the end.
Synthesis needs the findings, and the report needs the synthesis.

B. **[INCORRECT]** Run all four steps strictly one after another, so every step sees the output of the one before.
Search and analysis don't depend on each other, so waiting wastes time.

C. **[CORRECT]** Run web search and document analysis in parallel, then run synthesis once both complete, then report writing.
Correct. Independent work is parallelized and dependent work follows.

D. **[INCORRECT]** Run synthesis first so that it can tell search and analysis what to look for.
Synthesis depends on the findings from search and analysis.

Key concept: Parallel versus sequential orchestration (Task 1.2, 1.3).
Exam reasoning: Dependencies, not convenience, decide the order.
Production lesson: Parallelize the independent, sequence the dependent.

Q20 | Domain 1: Agentic Architecture & Orchestration | Scenario 3: Multi-Agent Research System

Task statement 1.3 | Intermediate | Single answer

Correct answer: D - Pass the batch 1 results (entity names and mappings) explicitly in the batch 2 prompt, since subagents don't share memory between invocations.

Why the correct answer is correct: Subagents do not automatically share memory between invocations; needed context must be explicitly provided in the prompt.

Option-by-option:

A. **[INCORRECT]** Run batch 2 before batch 1 so that the more complex documents are processed first.
Order doesn't create shared memory.

B. **[INCORRECT]** Tell the subagent in the batch 2 prompt to remember what it learned during batch 1 and to keep the entity names consistent with its earlier work, whatever they were.
It has no such memory to draw on.

C. **[INCORRECT]** Use a larger model for batch 2 because it can better infer earlier decisions.
Inference can't recover information that isn't present.

D. **[CORRECT]** Pass the batch 1 results (entity names and mappings) explicitly in the batch 2 prompt, since subagents don't share memory between invocations.
Correct. Each invocation starts without memory of previous ones.

Key concept: No shared memory between invocations (Task 1.3).
Exam reasoning: Contradicting names are consistent with two independent runs.
Production lesson: Never rely on a subagent remembering.

Q21 | Domain 1: Agentic Architecture & Orchestration | Scenario 3: Multi-Agent Research System

Task statement 1.7 | Intermediate | Single answer

Correct answer: D - Start a new session with a structured summary of the research plan and verified background, and re-run the searches for time-sensitive data.

Why the correct answer is correct: Starting a new session with a structured summary is more reliable than resuming with stale tool results.

Option-by-option:

A. **[INCORRECT]** Resume the old session and tell the model that some prices may have changed, then use its best judgment.
The model can't know which values changed.

B. **[INCORRECT]** Fork the old session and ask the fork to trust the earlier tool results.
A fork inherits the same stale results.

C. **[INCORRECT]** Resume the old session and ask for a refreshed report using the data that is already in its history, which saves the cost of repeating the searches.

The history contains outdated prices.

D. **[CORRECT]** Start a new session with a structured summary of the research plan and verified background, and re-run the searches for time-sensitive data.
Correct. Stale tool results shouldn't be carried forward.

Key concept: Stale results versus fresh start (Task 1.7).
Exam reasoning: Time-sensitive data has expired.
Production lesson: Re-fetch what ages.

Q22 | Domain 2: Tool Design & MCP Integration | Scenario 3: Multi-Agent Research System

Task statement 2.2 | Challenging | Single answer, secondary: D5

Correct answer: D - Return a permission-category, non-retryable error naming the cause; the subagent seeks an open-access version or another source, and reports a gap if none exists.

Why the correct answer is correct: Structured errors distinguish permission failures from transient ones and include retryability, enabling alternatives and honest coverage reporting.

Option-by-option:

A. **[INCORRECT]** Return an empty successful result so the subagent moves on without any message.
That hides the gap and implies nothing was found.

B. **[INCORRECT]** Return a transient error with isRetryable true so that the subagent keeps trying with increasing delays until the paywall eventually disappears or the access token is refreshed.
A paywall won't resolve on retry.

C. **[INCORRECT]** Return the HTML of the paywall page as the article content so the subagent can try to read it.
The paywall page isn't the article and would pollute the findings.

D. **[CORRECT]** Return a permission-category, non-retryable error naming the cause; the subagent seeks an open-access version or another source, and reports a gap if none exists.
Correct. The error guides recovery and avoids wasted retries.

Key concept: Permission errors and alternatives (Task 2.2, 5.3).
Exam reasoning: A 403 is not temporary, so retries are futile.
Production lesson: Help the agent choose a different path.

Q23 | Domain 2: Tool Design & MCP Integration | Scenario 3: Multi-Agent Research System

Task statement 2.4 | Intermediate | Single answer

Correct answer: C - Expand the MCP tool's description to explain what it covers, what it returns, and when it is preferable to a general web search.

Why the correct answer is correct: Enhance MCP tool descriptions to explain capabilities and outputs in detail so the agent doesn't prefer built-in tools over more capable MCP tools.

Option-by-option:

A. **[INCORRECT]** Remove WebSearch from all agents, so the licensed tool is the only search option anywhere in the system and no agent can be tempted by a built-in tool.
This breaks legitimate web research.

B. **[INCORRECT]** Add a rule in the system prompt: 'never use WebSearch under any circumstances'.
An absolute rule overrides judgment for tasks that need the web.

C. **[CORRECT]** Expand the MCP tool's description to explain what it covers, what it returns, and when it is preferable to a general web search.
Correct. A richer description lets the agent recognize the better tool.

D. **[INCORRECT]** Rename the tool to something longer and hope that its name attracts the model.
Names alone don't convey coverage or advantages.

Key concept: MCP tool descriptions (Task 2.4, 2.1).
Exam reasoning: A one-line description can't compete with a familiar built-in tool.
Production lesson: Sell the tool to the model in its description.

Q24 | Domain 2: Tool Design & MCP Integration | Scenario 3: Multi-Agent Research System

Task statement 2.3 | Intermediate | Single answer, secondary: D1

Correct answer: D - Restrict the coordinator to a small set of orchestration tools, chiefly Task, and leave the specialist tools with the subagents.

Why the correct answer is correct: Agents perform better with scoped tools; coordinators should hold orchestration tools while specialists hold domain tools.

Option-by-option:

A. **[INCORRECT]** Keep the tool set but add a paragraph to the prompt telling the coordinator not to use specialist tools directly and to delegate such work.
Prompt restraint is probabilistic, and the large tool set remains.

B. **[INCORRECT]** Give the subagents the same 20 extra tools so that everyone has the same capabilities.
This worsens selection for every agent.

C. **[INCORRECT]** Force tool_choice to the Task tool on every turn so that the coordinator always delegates.
Forcing Task every turn prevents the coordinator from ever finishing or synthesizing.

D. **[CORRECT]** Restrict the coordinator to a small set of orchestration tools, chiefly Task, and leave the specialist tools with the subagents.
Correct. Role-specific tool sets keep delegation clean and selection reliable.

Key concept: Tool distribution across agents (Task 2.3, 1.2).
Exam reasoning: The coordinator's tools duplicate its subagents', inviting it to bypass them.
Production lesson: Tool access defines behavior.

Q25 | Domain 2: Tool Design & MCP Integration | Scenario 3: Multi-Agent Research System

Task statement 2.1 | Foundational | Select 2

Correct answer: A, D - Input formats and example queries that illustrate typical use. / Boundary statements describing when to use this tool versus a similar one.

Why the correct answer is correct: Tool descriptions should include input formats, example queries, edge cases, and boundary explanations.

Option-by-option:

A. **[CORRECT]** Input formats and example queries that illustrate typical use.
Correct. Examples and formats help the model match requests to the tool.

B. **[INCORRECT]** Enthusiastic wording stating that the tool is fast, powerful, and reliable.
Promotional language adds no selection information.

C. **[INCORRECT]** The tool's internal ticket number and the date it was last deployed.
Administrative metadata doesn't help choose a tool.

D. **[CORRECT]** Boundary statements describing when to use this tool versus a similar one.
Correct. Boundaries prevent misrouting.

E. **[INCORRECT]** A reduction of the description to under five words for readability.
Minimal descriptions cause unreliable selection.

Key concept: Description content (Task 2.1).
Exam reasoning: Two options add selection information; three don't.
Production lesson: Describe usage, not marketing.

Q26 | Domain 2: Tool Design & MCP Integration | Scenario 3: Multi-Agent Research System

Task statement 2.1 | Intermediate | Single answer

Correct answer: B - Rename them to reflect their roles (for example search_internal_kb and fetch_web_page) and rewrite each description with its source, inputs, outputs, and use.

Why the correct answer is correct: Rename tools and update descriptions to eliminate functional overlap; clear differentiation of purpose, inputs, and outputs fixes misrouting.

Option-by-option:

A. **[INCORRECT]** Combine the two into a single get_data tool with a source argument.
This hides two different capabilities behind a vague interface.

B. **[CORRECT]** Rename them to reflect their roles (for example search_internal_kb and fetch_web_page) and rewrite each description with its source, inputs, outputs, and use.
Correct. Distinct names and descriptions eliminate the overlap.

C. **[INCORRECT]** Add a third tool that decides which of the two to call and then calls it.
A router tool adds complexity without fixing the descriptions.

D. **[INCORRECT]** Keep the names and add a one-line prompt: 'prefer get_info for questions about the company and fetch_data for questions about the outside world' whenever possible.
A preference rule doesn't clarify the purpose and might misroute web fetches.

Key concept: Eliminating overlap (Task 2.1).
Exam reasoning: The tools do different things but look the same.
Production lesson: Make their differences visible.

Q27 | Domain 4: Prompt Engineering & Structured Output | Scenario 3: Multi-Agent Research System

Task statement 4.6 | Foundational | Select 3

Correct answer: A, B, F - A model that generated content retains its reasoning context, which makes it less likely to question its own decisions in the same session. / Splitting a large review into per-file passes plus a cross-file integration pass helps avoid attention dilution and contradictory findings. / An independent instance without the generator's reasoning context is more effective at catching subtle issues than self-review instructions or extended thinking.

Why the correct answer is correct: Self-review is limited by retained reasoning context; independent instances and multi-pass splits improve review quality.

Option-by-option:

A. **[CORRECT]** A model that generated content retains its reasoning context, which makes it less likely to question its own decisions in the same session.
Correct. This is the self-review limitation.

B. **[CORRECT]** Splitting a large review into per-file passes plus a cross-file integration pass helps avoid attention dilution and contradictory findings.
Correct. Focused passes improve consistency.

C. **[INCORRECT]** A larger context window eliminates attention dilution, so a large review no longer needs to be split.
Window size does not fix attention quality.

D. **[INCORRECT]** Telling the same session to act as a skeptical reviewer is equivalent to using a separate instance.
The retained context still biases the review.

E. **[INCORRECT]** Extended thinking removes the need for an independent reviewer on important changes.
It does not provide independence from the generator's context.

F. **[CORRECT]** An independent instance without the generator's reasoning context is more effective at catching subtle issues than self-review instructions or extended thinking.
Correct. Independence matters.

Key concept: Multi-instance and multi-pass reviews (Task 4.6).
Exam reasoning: The correct statements are the three documented findings; the wrong ones are common misconceptions.
Production lesson: Independence and focus beat instruction.

Q28 | Domain 5: Context Management & Reliability | Scenario 3: Multi-Agent Research System

Task statement 5.3 | Challenging | Single answer

Correct answer: C - Use separate sections for well-established findings, contested findings with each source's position, and a coverage section stating gaps from unavailable sources.

Why the correct answer is correct: Synthesis output should include coverage annotations indicating which findings are well supported and which areas have gaps, and should distinguish well-established from contested findings.

Option-by-option:

A. **[INCORRECT]** Remove contested findings and gaps from the report to keep it clean and authoritative.
Omitting them overstates certainty.

B. **[INCORRECT]** Keep the single list but order the bullets by importance, trusting the reader to infer the level of confidence from each bullet's position in the list.
Order doesn't express evidence strength or gaps.

C. **[CORRECT]** Use separate sections for well-established findings, contested findings with each source's position, and a coverage section stating gaps from unavailable sources.
Correct. Structure reveals confidence and gaps.

D. **[INCORRECT]** Add a footer note saying 'some findings may be uncertain' without identifying which ones.
A generic caveat is not actionable.

Key concept: Coverage annotations (Task 5.3, 5.6).
Exam reasoning: The readers' inability to separate strong from weak is the defect.
Production lesson: Show the shape of the evidence.

Q29 | Domain 5: Context Management & Reliability | Scenario 3: Multi-Agent Research System

Task statement 5.3 | Intermediate | Single answer, secondary: D2

Correct answer: D - A successful result with zero matches, stating the queries and sources searched, so the coordinator can report 'no studies found' accurately.

Why the correct answer is correct: Distinguish access failures from valid empty results; a successful query with no matches is a legitimate result and should be reported with what was searched.

Option-by-option:

A. **[INCORRECT]** Nothing at all, since an empty result has no value to the coordinator.
The absence of studies after a documented search is itself a finding.

B. **[INCORRECT]** A failure with category transient and isRetryable true, so that the coordinator retries the search repeatedly for this topic until something is found.
Nothing failed, and retries won't create studies.

C. **[INCORRECT]** A fabricated summary of related studies to avoid reporting an empty result.
Fabrication is unacceptable.

D. **[CORRECT]** A successful result with zero matches, stating the queries and sources searched, so the coordinator can report 'no studies found' accurately.
Correct. A completed search with no matches is a valid, informative result.

Key concept: Valid empty results (Task 5.3, 2.2).
Exam reasoning: The search succeeded, so it should not be labeled a failure.
Production lesson: Report absence honestly and with its method.

Q30 | Domain 1: Agentic Architecture & Orchestration | Scenario 5: Claude Code for Continuous Integration

Task statement 1.6 | Intermediate | Single answer, secondary: D3

Correct answer: B - A chain: analyze new code and existing tests for gaps, propose uncovered scenarios, generate tests for those, then run them and fix failures.

Why the correct answer is correct: Prompt chaining breaks tasks into focused sequential steps; test generation benefits from knowing existing tests and gaps.

Option-by-option:

A. **[INCORRECT]** Generate tests only for functions that changed, ignoring what the test suite already contains.
Ignoring existing tests causes duplication.

B. **[CORRECT]** A chain: analyze new code and existing tests for gaps, propose uncovered scenarios, generate tests for those, then run them and fix failures.
Correct. Each step builds on the previous one and focuses on one task.

C. **[INCORRECT]** A single request telling Claude to generate as many tests as possible for every function it can find, so that coverage is as high as possible.
Volume produces duplicates and low-value tests.

D. **[INCORRECT]** Generate tests twice with different seeds and merge both sets into the repository.
Merging two outputs doubles duplication.

Key concept: Prompt chaining for test generation (Task 1.6, 3.6).
Exam reasoning: Duplication and missed branches both come from not analyzing what exists first.
Production lesson: Look before you generate.

Q31 | Domain 1: Agentic Architecture & Orchestration | Scenario 5: Claude Code for Continuous Integration

Task statement 1.6 | Challenging | Single answer, secondary: D2

Correct answer: D - Have the integration pass find callers of the changed interface (for example with Grep) and check those call sites against the new contract, even outside the diff.

Why the correct answer is correct: A separate integration pass examines cross-file data flow; for interface changes this must include dependent code outside the diff.

Option-by-option:

A. **[INCORRECT]** Add an instruction to each per-file pass to speculate about which other files might be affected.
Speculation without looking at the call sites is unreliable.

B. **[INCORRECT]** Give the integration pass only the diff and ask it to think harder about possible side effects on the rest of the codebase and to list them in detail for the developer.
The affected code isn't in the diff, so it cannot be evaluated.

C. **[INCORRECT]** Rely on the unit tests to catch it and skip any cross-file analysis in the review.
Tests may not cover the call site, and the review is meant to catch such issues.

D. **[CORRECT]** Have the integration pass find callers of the changed interface (for example with Grep) and check those call sites against the new contract, even outside the diff.
Correct. Cross-file breakage requires examining unchanged dependents.

Key concept: Cross-file integration review (Task 1.6, 2.5).
Exam reasoning: The broken caller is outside the diff.
Production lesson: Integration review must follow the dependencies.

Q32 | Domain 1: Agentic Architecture & Orchestration | Scenario 5: Claude Code for Continuous Integration

Task statement 1.6 | Challenging | Single answer

Correct answer: A - Adapt the plan: prioritize creating tests for the 'money' module first, then return to ledger, updating the plan as dependencies are discovered.

Why the correct answer is correct: Open-ended tasks benefit from adaptive plans: map structure, prioritize high-impact areas, and adjust as dependencies are discovered.

Option-by-option:

A. **[CORRECT]** Adapt the plan: prioritize creating tests for the 'money' module first, then return to ledger, updating the plan as dependencies are discovered.
Correct. Adaptive decomposition responds to discovered dependencies.

B. **[INCORRECT]** Restart the nightly job from scratch each time a dependency is found.
Restarting wastes completed work.

C. **[INCORRECT]** Continue in alphabetical order and accept the failing ledger tests until the 'money' module is eventually reached, since the plan was agreed in advance.
A fixed order produces known-broken results.

D. **[INCORRECT]** Skip any module that has failing tests and report them as untestable.
The modules are testable once the dependency is handled.

Key concept: Adaptive decomposition (Task 1.6).
Exam reasoning: The new dependency information should change the order.
Production lesson: Plans should respond to what you learn.

Q33 | Domain 2: Tool Design & MCP Integration | Scenario 5: Claude Code for Continuous Integration

Task statement 2.2 | Challenging | Single answer, secondary: D3

Correct answer: B - Return a permission-category, non-retryable error with a clear description; the agent stops retrying and outputs the findings to the job summary instead.

Why the correct answer is correct: Structured errors with categories and retryability prevent wasted retries, and non-retryable permission errors should lead to an alternative path.

Option-by-option:

A. **[INCORRECT]** Return a transient error so the agent keeps retrying with backoff until the token is refreshed by an administrator or the job times out, whichever happens first.
The permission problem won't resolve itself during the job.

B. **[CORRECT]** Return a permission-category, non-retryable error with a clear description; the agent stops retrying and outputs the findings to the job summary instead.
Correct. The structured error prevents futile retries and the findings are still delivered.

C. **[INCORRECT]** Return success without posting, so that the CI check passes and nobody is blocked.
The comments were never delivered, and the pipeline would hide that.

D. **[INCORRECT]** Return the stack trace of the HTTP client so that the model can work out the cause.
Unstructured traces are noisy and don't tell the agent whether to retry.

Key concept: Permission errors in CI (Task 2.2, 3.6).
Exam reasoning: The failure isn't temporary, so retries are futile.
Production lesson: Make the error say 'stop' and 'try this instead'.

Q34 | Domain 2: Tool Design & MCP Integration | Scenario 5: Claude Code for Continuous Integration

Task statement 2.4 | Foundational | Single answer

Correct answer: D - Reference an environment variable, such as ${JIRA_TOKEN}, and provide the value from the CI system's secret store.

Why the correct answer is correct: Project-scoped .mcp.json supports environment variable expansion (for example ${GITHUB_TOKEN}) for credential management without committing secrets.

Option-by-option:

A. **[INCORRECT]** Add the token to the project's CLAUDE.md so the agent can read it when needed.
CLAUDE.md is committed and would expose the secret.

B. **[INCORRECT]** Ask the model to prompt a developer for the token each time the job runs.
CI is non-interactive and cannot prompt.

C. **[INCORRECT]** Write the token as a literal string in .mcp.json, since the repository is private.
A literal secret in a file is exposed to everyone with repository access.

D. **[CORRECT]** Reference an environment variable, such as ${JIRA_TOKEN}, and provide the value from the CI system's secret store.
Correct. Environment variable expansion keeps secrets out of version control.

Key concept: MCP credentials (Task 2.4).
Exam reasoning: The requirement not to commit secrets points to expansion.
Production lesson: Secrets live in the environment, not in the repo.

Q35 | Domain 3: Claude Code Configuration & Workflows | Scenario 5: Claude Code for Continuous Integration

Task statement 3.6 | Challenging | Single answer

Correct answer: C - Run in non-interactive mode with -p, request --output-format json with a --json-schema for findings, and document review criteria in CLAUDE.md.

Why the correct answer is correct: -p runs Claude Code non-interactively; --output-format json with --json-schema yields machine-parseable findings; CLAUDE.md provides project context such as review criteria.

Option-by-option:

A. **[INCORRECT]** Run with -p and ask in the prompt for the output in a particular markdown layout, without any CLAUDE.md.
Markdown layout drifts, and criteria are not persisted.

B. **[INCORRECT]** Run interactively with a here-document supplying answers to any prompts, and parse the text with regular expressions.
This is fragile and doesn't give structured output.

C. **[CORRECT]** Run in non-interactive mode with -p, request --output-format json with a --json-schema for findings, and document review criteria in CLAUDE.md.
Correct. Each element solves one of the problems.

D. **[INCORRECT]** Use --output-format json without -p and put the review criteria in the user's ~/.claude/CLAUDE.md.
Without -p the job may still hang, and CI doesn't read a developer's personal file.

Key concept: CI invocation (Task 3.6).
Exam reasoning: Three symptoms map to three features.
Production lesson: Know the CI trio: -p, json/schema, CLAUDE.md.

Q36 | Domain 3: Claude Code Configuration & Workflows | Scenario 5: Claude Code for Continuous Integration

Task statement 3.6 | Challenging | Single answer, secondary: D4

Correct answer: B - Run it on a labeled set of past PRs with known defects and known-clean changes, and measure precision and recall, including per category.

Why the correct answer is correct: Evaluate against known defects to measure false positives and false negatives; this supports decisions on blocking versus advisory. (Evaluation harness design is supplemental practice relating to CI review quality.)

Option-by-option:

A. **[INCORRECT]** Ask developers whether the comments seem helpful on their next ten PRs and decide by general impression, then switch on blocking if most say yes.
Impressions miss missed defects and are subjective.

B. **[CORRECT]** Run it on a labeled set of past PRs with known defects and known-clean changes, and measure precision and recall, including per category.
Correct. Known defects allow measurement of both error types.

C. **[INCORRECT]** Count how many comments Claude posts per PR and treat higher counts as better coverage.
Volume does not equal accuracy.

D. **[INCORRECT]** Run the review twice on each PR and measure only whether the two runs agree.
Agreement doesn't show correctness. Both runs can be wrong.

Note: this item draws partly on **supplemental** engineering knowledge that is not a confirmed exam objective.
Key concept: Evaluation of CI review quality (Task 3.6, 4.1; supplemental practice).
Exam reasoning: Only a labeled set shows what is missed as well as what is wrongly flagged.
Production lesson: Measure before you gate.

Q37 | Domain 3: Claude Code Configuration & Workflows | Scenario 5: Claude Code for Continuous Integration

Task statement 3.6 | Intermediate | Single answer

Correct answer: A - Keep automated build, test, and review on every PR, but add a manual approval gate before production deployment.

Why the correct answer is correct: CI automates integration, build, and test; continuous delivery keeps every change releasable with a manual release decision; continuous deployment releases automatically. (General CI/CD knowledge; supplemental to the exam guide.)

Option-by-option:

A. **[CORRECT]** Keep automated build, test, and review on every PR, but add a manual approval gate before production deployment.
Correct. Automated deployment to production without approval is continuous deployment.

B. **[INCORRECT]** Make the Claude review the approval gate that deploys to production when it finds no issues.
A model review is not a substitute for human approval of releases.

C. **[INCORRECT]** Keep everything as is, because deploying automatically after tests pass is what continuous integration means.
Automatic production release is continuous deployment, not CI.

D. **[INCORRECT]** Remove automated testing from the pipeline, since CI is only about merging code frequently and tests slow down the integration of every change into the main branch.
CI includes automated build and test.

Note: this item draws partly on **supplemental** engineering knowledge that is not a confirmed exam objective.
Key concept: CI versus delivery versus deployment (Task 3.6; supplemental).
Exam reasoning: The stated goal of human-approved releases excludes automatic production deployment.
Production lesson: Know where the human gate sits.

Q38 | Domain 3: Claude Code Configuration & Workflows | Scenario 5: Claude Code for Continuous Integration

Task statement 3.6 | Intermediate | Single answer

Correct answer: A - Use a token limited to reading code and posting PR comments, and keep branch and settings write access out of that job.

Why the correct answer is correct: Apply least privilege to automated agents: grant only the permissions the role requires. (Security practice; supplemental to the exam guide.)

Option-by-option:

A. **[CORRECT]** Use a token limited to reading code and posting PR comments, and keep branch and settings write access out of that job.
Correct. Least privilege limits the damage of any mistake or injection.

B. **[INCORRECT]** Keep the broad token and add a prompt instruction saying that the job must never push branches or edit repository settings under any circumstances.
A prompt cannot technically prevent those actions.

C. **[INCORRECT]** Remove authentication from the job so that it fails if it tries to do anything sensitive.
Without authentication it also cannot read or comment.

D. **[INCORRECT]** Keep the broad token but log every action, and review the logs weekly.
Logs detect misuse after the fact.

Note: this item draws partly on **supplemental** engineering knowledge that is not a confirmed exam objective.
Key concept: Least privilege in CI (Task 3.6; supplemental).
Exam reasoning: The job's needs are narrow, so the credentials should be too.
Production lesson: Constrain by permission, not by promise.

Q39 | Domain 4: Prompt Engineering & Structured Output | Scenario 5: Claude Code for Continuous Integration

Task statement 4.6 | Challenging | Single answer, secondary: D3

Correct answer: A - Give the reviewer the diff, requirements, and project criteria, but not the generator's reasoning, so that it evaluates the change on its own.

Why the correct answer is correct: Independent review instances without prior reasoning context are more effective at catching subtle issues; passing the generator's reasoning defeats the independence.

Option-by-option:

A. **[CORRECT]** Give the reviewer the diff, requirements, and project criteria, but not the generator's reasoning, so that it evaluates the change on its own.
Correct. Independence requires withholding the generator's reasoning context.

B. **[INCORRECT]** Give the reviewer even more of the generator's reasoning and commentary, so that it understands the intent fully before judging the change and avoids false alarms.
More reasoning increases anchoring on the generator's view.

C. **[INCORRECT]** Use the same session as the generator for the review, with a prompt telling it to be strict.
The same session retains the reasoning that it was supposed to question.

D. **[INCORRECT]** Ask the reviewer to disagree with the generator at least twice per review.
Forced disagreement produces noise rather than insight.

Key concept: Review independence (Task 4.6, 3.6).
Exam reasoning: Agreement with the generator results from sharing its reasoning.
Production lesson: Independent means uninformed of the author's logic.

Q40 | Domain 4: Prompt Engineering & Structured Output | Scenario 5: Claude Code for Continuous Integration

Task statement 4.6 | Intermediate | Single answer

Correct answer: B - Give the review explicit criteria for what to report and skip and use focused per-file passes, instead of requiring agreement among runs.

Why the correct answer is correct: Consensus voting can suppress real bugs that are only intermittently detected; explicit criteria and multi-pass structure improve precision more directly.

Option-by-option:

A. **[INCORRECT]** Keep the voting and lower the severity of all findings to hide the missed ones.
Lowering severity doesn't recover missed bugs.

B. **[CORRECT]** Give the review explicit criteria for what to report and skip and use focused per-file passes, instead of requiring agreement among runs.
Correct. It addresses the cause of noise without suppressing intermittent true findings.

C. **[INCORRECT]** Require agreement in all three runs instead of two, and drop single-run findings entirely, for even stronger filtering of the noisy findings that developers complain about.
This suppresses even more true findings.

D. **[INCORRECT]** Run the review five times and post the union of all findings.
The union increases noise.

Key concept: Precision versus recall trade-offs (Task 4.6, 4.1).
Exam reasoning: Real bugs missed because they appear in one run is the cost of voting.
Production lesson: Fix the cause of noise, not the symptom.

Q41 | Domain 4: Prompt Engineering & Structured Output | Scenario 5: Claude Code for Continuous Integration

Task statement 4.1 | Challenging | Single answer

Correct answer: A - State the categories to report (bugs, security issues, correctness risks) and the categories to skip (minor style, locally accepted patterns), not confidence.

Why the correct answer is correct: General instructions like 'be conservative' or 'only report high-confidence findings' fail to improve precision compared with specific categorical criteria.

Option-by-option:

A. **[CORRECT]** State the categories to report (bugs, security issues, correctness risks) and the categories to skip (minor style, locally accepted patterns), not confidence.
Correct. Specific categorical criteria outperform confidence filtering.

B. **[INCORRECT]** Add the instruction 'be conservative' to the end of the prompt.
General instructions don't improve precision.

C. **[INCORRECT]** Strengthen the confidence instruction to 'only report findings you are 95% sure about, and say nothing about anything less certain than that in the diff'.
A numeric confidence demand doesn't define what counts as a real problem.

D. **[INCORRECT]** Ask the model to rank all findings and post only the top three per PR.
Ranking doesn't remove style noise and may drop real issues.

Key concept: Explicit criteria (Task 4.1).
Exam reasoning: The noise is a category problem, not a confidence problem.
Production lesson: Define in or out, not how sure.

Q42 | Domain 4: Prompt Engineering & Structured Output | Scenario 5: Claude Code for Continuous Integration

Task statement 4.6 | Intermediate | Single answer

Correct answer: C - Per-file local passes, then a cross-file integration pass, then an independent verification pass on any finding that would block the merge.

Why the correct answer is correct: Multi-pass reviews split local and integration analysis, and independent instances verify high-impact findings.

Option-by-option:

A. **[INCORRECT]** Per-file passes only, with all findings posted directly as blocking comments.
No cross-file analysis, and unverified findings block merges.

B. **[INCORRECT]** One pass over the whole PR, followed by a second pass by the same session to double-check and refine all of its own findings before posting.
Single-pass attention dilution remains, and self-review is weak.

C. **[CORRECT]** Per-file local passes, then a cross-file integration pass, then an independent verification pass on any finding that would block the merge.
Correct. Depth, integration, and independent verification of high-impact findings.

D. **[INCORRECT]** Three parallel full-PR passes with majority voting and no integration pass.
Voting can suppress true findings and nothing targets cross-file flows.

Key concept: Multi-pass architecture (Task 4.6).
Exam reasoning: Only one option combines all three needed elements.
Production lesson: Layer the review.

Q43 | Domain 5: Context Management & Reliability | Scenario 5: Claude Code for Continuous Integration

Task statement 5.5 | Foundational | Single answer

Correct answer: A - An overall figure can hide poor accuracy in particular categories, so precision should be measured per finding category before choosing which, if any, to block on.

Why the correct answer is correct: Aggregate accuracy metrics can mask poor performance on specific segments; validate accuracy by category before relying on it.

Option-by-option:

A. **[CORRECT]** An overall figure can hide poor accuracy in particular categories, so precision should be measured per finding category before choosing which, if any, to block on.
Correct. Aggregate accuracy may mask weak segments.

B. **[INCORRECT]** The figure is sufficient if the review sample is large, because a big sample makes per-category breakdowns unnecessary.
A large sample still does not show whether particular categories perform worse than the average.

C. **[INCORRECT]** Blocking depends only on how fast the review runs, not on its accuracy.
Accuracy matters directly for blocking decisions.

D. **[INCORRECT]** The figure is insufficient only if it was computed by developers instead of by the model.
Who computes it is not the issue.

Key concept: Segmented accuracy (Task 5.5).
Exam reasoning: A single average hides category differences.

Production lesson: Break the average apart.

Q44 | Domain 1: Agentic Architecture & Orchestration | Scenario 6: Structured Data Extraction

Task statement 1.5 | Intermediate | Single answer

Correct answer: B - A pre-call hook on the write tool that records the record hash, source document ID, and run ID to an audit log before allowing the call.

Why the correct answer is correct: Hooks that intercept outgoing tool calls provide deterministic enforcement and instrumentation, independent of model compliance.

Option-by-option:

A. **[INCORRECT]** A prompt instruction asking the agent to include a source document ID and a run ID in the free-text note of every write it makes to the ledger.
The note may be omitted or wrong, and it isn't a durable audit record.

B. **[CORRECT]** A pre-call hook on the write tool that records the record hash, source document ID, and run ID to an audit log before allowing the call.
Correct. The hook executes on every write, so logging cannot be skipped.

C. **[INCORRECT]** A weekly script that reconstructs which document produced each ledger entry by matching amounts.
Matching amounts is ambiguous and happens after the fact.

D. **[INCORRECT]** An optional field in the schema where the model may add a reference to the source if it chooses.
Optional model-supplied fields aren't guaranteed.

Key concept: Hooks for guaranteed behavior (Task 1.5).
Exam reasoning: 'Regardless of whether the model remembers' requires code-level enforcement.
Production lesson: Put audit obligations in code paths.

Q45 | Domain 1: Agentic Architecture & Orchestration | Scenario 6: Structured Data Extraction

Task statement 1.5 | Foundational | Single answer

Correct answer: A - A PostToolUse hook that maps numeric codes to labeled statuses before the model sees the result.

Why the correct answer is correct: PostToolUse hooks can normalize numeric status codes and other heterogeneous formats before the model processes tool results.

Option-by-option:

A. **[CORRECT]** A PostToolUse hook that maps numeric codes to labeled statuses before the model sees the result.
Correct. Hooks normalize results before the model processes them.

B. **[INCORRECT]** A request to the tool's owners to renumber the codes so that failures are small numbers.
That depends on others and doesn't fix existing consumers.

C. **[INCORRECT]** A system prompt paragraph that says higher status codes are worse and should be treated as failures.
The rule is easy to misapply and needs the model to translate each time.

D. **[INCORRECT]** An example showing a status code 3 being handled correctly, placed in the extraction prompt.
One example doesn't guarantee consistent interpretation.

Key concept: PostToolUse normalization (Task 1.5).
Exam reasoning: The model needs labels, not codes.
Production lesson: Translate at the boundary.

Q46 | Domain 1: Agentic Architecture & Orchestration | Scenario 6: Structured Data Extraction

Task statement 1.5 | Challenging | Single answer, secondary: D5

Correct answer: A - A circuit-breaker hook that halts automated posting and redirects the batch to review when the validation failure rate crosses a set threshold.

Why the correct answer is correct: Hooks and gates can enforce policy deterministically at system level, including halting and redirecting when error rates indicate systemic failure.

Option-by-option:

A. **[CORRECT]** A circuit-breaker hook that halts automated posting and redirects the batch to review when the validation failure rate crosses a set threshold.
Correct. It stops a systemic problem from propagating and routes to humans.

B. **[INCORRECT]** A rule that retries each failed document indefinitely until it validates.
Unbounded retries may never succeed and could fabricate values.

C. **[INCORRECT]** A nightly cleanup job that removes bad ledger entries after the batch finishes.
Cleanup after posting leaves bad data in the system in the meantime.

D. **[INCORRECT]** A prompt instruction telling the agent to slow down and become more careful when it notices that many of the documents it processes fail validation.
The agent doesn't see batch-level statistics and prompts aren't guarantees.

Key concept: Deterministic guardrails for batch processing (Task 1.5, 5.5).
Exam reasoning: The failure is systemic, so the guard must observe the batch, not a single document.
Production lesson: Stop the line when quality collapses.

Q47 | Domain 2: Tool Design & MCP Integration | Scenario 6: Structured Data Extraction

Task statement 2.3 | Intermediate | Single answer

Correct answer: C - Give the agent a narrow get_fx_rate tool for the one lookup it needs, and keep the rest of the finance tools out of its tool set.

Why the correct answer is correct: Provide scoped cross-role tools for specific needs instead of broad tool access, preserving least privilege and selection reliability.

Option-by-option:

A. **[INCORRECT]** Connect the agent to all 17 finance tools so that it can handle any financial need it encounters without asking for new access later.
More tools degrade selection and expose payment operations.

B. **[INCORRECT]** Ask the agent to estimate exchange rates from its general knowledge, which is usually close.
Estimated rates produce incorrect monetary values.

C. **[CORRECT]** Give the agent a narrow get_fx_rate tool for the one lookup it needs, and keep the rest of the finance tools out of its tool set.
Correct. A scoped tool meets the need without exposing sensitive operations.

D. **[INCORRECT]** Pass the rate table in every prompt regardless of whether a conversion is needed.
This wastes context and may be stale.

Key concept: Scoped tool access (Task 2.3).
Exam reasoning: The need is one lookup; the exposure would be seventeen tools.
Production lesson: Expose the minimum.

Q48 | Domain 2: Tool Design & MCP Integration | Scenario 6: Structured Data Extraction

Task statement 2.3 | Intermediate | Single answer, secondary: D1

Correct answer: C - Remove post_to_ledger from the agent's tools, and have deterministic pipeline code post only records that have passed validation.

Why the correct answer is correct: Restrict each agent to the tools needed for its role; side-effect actions can be moved to deterministic code after validation.

Option-by-option:

A. **[INCORRECT]** Keep post_to_ledger but set tool_choice to 'auto' so the agent uses it only when needed.
'auto' leaves the decision to the model.

B. **[INCORRECT]** Keep post_to_ledger and rename it submit_for_review so that the agent treats it more carefully.
Renaming doesn't change what the tool does.

C. **[CORRECT]** Remove post_to_ledger from the agent's tools, and have deterministic pipeline code post only records that have passed validation.
Correct. The agent cannot call what it does not have, and posting follows verified checks.

D. **[INCORRECT]** Keep post_to_ledger and instruct the agent to call it only after the user confirms the extraction, as the final step of each document.
Users aren't present in batch processing, and prompt compliance isn't guaranteed.

Key concept: Least privilege for side effects (Task 2.3, 1.4).
Exam reasoning: The failure was an agent with a capability that it should not have had.
Production lesson: Separate thinking from doing.

Q49 | Domain 2: Tool Design & MCP Integration | Scenario 6: Structured Data Extraction

Task statement 2.2 | Intermediate | Single answer

Correct answer: D - Return an isError result with category validation, isRetryable false, and the field-level errors, so the agent or pipeline can correct the input.

Why the correct answer is correct: Structured errors with category, retryability, and descriptions let agents choose appropriate recovery; validation errors require correcting input.

Option-by-option:

A. **[INCORRECT]** Swallow the error and report success to avoid blocking the pipeline.
Hiding failure leads to missing ledger entries.

B. **[INCORRECT]** Return 'Posting failed' with a hint that the agent should try again after a short wait.
Waiting doesn't fix an invalid cost center.

C. **[INCORRECT]** Return a transient error so that the retry logic handles it automatically, since most 4xx errors clear after a while and the posting will then go through.
Validation errors are not transient.

D. **[CORRECT]** Return an isError result with category validation, isRetryable false, and the field-level errors, so the agent or pipeline can correct the input.
Correct. The details allow a corrected resubmission rather than blind retries.

Key concept: Validation errors (Task 2.2).
Exam reasoning: A 422 with field errors is precisely actionable information.
Production lesson: Pass useful detail through.

Q50 | Domain 2: Tool Design & MCP Integration | Scenario 6: Structured Data Extraction

Task statement 2.1 | Intermediate | Single answer, secondary: D4

Correct answer: C - State in the description and schema that amounts are integers in minor currency units (cents) alongside a separate ISO currency code field.

Why the correct answer is correct: Tool descriptions and schemas should specify outputs precisely, including units and formats, to prevent misinterpretation by the model.

Option-by-option:

A. **[INCORRECT]** Return the amount as a formatted string such as '$1,250.00' for readability.
Formatted strings are harder to validate and may lose the currency code.

B. **[INCORRECT]** Leave the contract unchanged and ask downstream agents to infer the unit from magnitude.
Inference is unreliable for small and large amounts.

C. **[CORRECT]** State in the description and schema that amounts are integers in minor currency units (cents) alongside a separate ISO currency code field.
Correct. Precise output contracts prevent misinterpretation.

D. **[INCORRECT]** Add a sentence to the description saying that the amount is 'accurate', 'verified', and 'should be used carefully' by every downstream agent that reads it.
Vague assurances do not define units.

Key concept: Precise tool contracts (Task 2.1, 4.3).
Exam reasoning: The inconsistency stems from an undefined unit.
Production lesson: Define units in the contract.

Q51 | Domain 4: Prompt Engineering & Structured Output | Scenario 6: Structured Data Extraction

Task statement 4.4 | Intermediate | Single answer, secondary: D5

Correct answer: B - Limit retries to a small number (for example two), then route the document to human review with the document, the attempts, and the errors attached.

Why the correct answer is correct: Retry with error feedback works for format and structural errors but is ineffective when information is absent; failing cases should be routed for human review.

Option-by-option:

A. **[INCORRECT]** Raise the limit to twenty rounds so that the model has more chances to get it right, with the validation error attached on every round, whatever the cost.
If the information is absent or the problem structural, more rounds don't help.

B. **[CORRECT]** Limit retries to a small number (for example two), then route the document to human review with the document, the attempts, and the errors attached.
Correct. Bounded retries plus escalation prevent endless loops.

C. **[INCORRECT]** Drop documents that fail validation twice, to keep the pipeline moving.
Dropping documents loses data without a record.

D. **[INCORRECT]** After the final failure, fill the failing fields with defaults and accept the record.
Defaults hide the problem and introduce incorrect data.

Key concept: Bounded retries (Task 4.4, 5.5).
Exam reasoning: Repeated failure signals a non-retryable cause.
Production lesson: Know when to stop retrying.

Q52 | Domain 4: Prompt Engineering & Structured Output | Scenario 6: Structured Data Extraction

Task statement 4.5 | Intermediate | Single answer

Correct answer: D - Restructure to a single-turn extraction per batch request, and do vendor lookups in pipeline code before or after the batch, since batch requests can't run tools mid-request.

Why the correct answer is correct: The Batches API does not support multi-turn tool calling within a single request; restructure so that tool-dependent steps happen outside the batch.

Option-by-option:

A. **[INCORRECT]** Switch the whole job to the synchronous API since any tool use rules out batch processing.
Single-turn structured extraction is fine in batch; only mid-request tool loops are a problem.

B. **[INCORRECT]** Embed the entire vendor database in each request to avoid needing a lookup.
That inflates every request's size and cost.

C. **[INCORRECT]** Submit the batch as is, and expect each request to call the lookup tool, receive the vendor data, and continue its extraction in the very same request, as it does today in the synchronous flow.
Batch requests can't execute tools mid-request and return results.

D. **[CORRECT]** Restructure to a single-turn extraction per batch request, and do vendor lookups in pipeline code before or after the batch, since batch requests can't run tools mid-request.
Correct. Batches don't support multi-turn tool calling within a request.

Key concept: Batch limitations (Task 4.5).
Exam reasoning: A mid-request tool loop conflicts with batch semantics.
Production lesson: Fit the workflow to the API's constraints.

Q53 | Domain 4: Prompt Engineering & Structured Output | Scenario 6: Structured Data Extraction

Task statement 4.3 | Intermediate | Single answer

Correct answer: C - Represent amounts as integers in minor currency units, or as decimal strings, with a separate currency code, and do arithmetic in decimal types.

Why the correct answer is correct: Schema design should avoid representations that lose precision for monetary values and include explicit currency information; semantic checks still apply.

Option-by-option:

A. **[INCORRECT]** Store amounts as free text and let each downstream consumer parse them as it prefers.
Inconsistent parsing causes new discrepancies.

B. **[INCORRECT]** Keep floats but tell the model in the prompt instructions to round every value to two decimals before it fills in the field, so the cents are right.
Rounding instructions don't remove floating-point arithmetic errors.

C. **[CORRECT]** Represent amounts as integers in minor currency units, or as decimal strings, with a separate currency code, and do arithmetic in decimal types.
Correct. Avoiding binary floating point preserves exact values.

D. **[INCORRECT]** Use floats but multiply by 1.0001 to offset the observed difference.
An arbitrary correction factor is wrong.

Key concept: Monetary schema design (Task 4.3).
Exam reasoning: A one-cent drift is the signature of floating-point representation.
Production lesson: Money is not a float.

Q54 | Domain 4: Prompt Engineering & Structured Output | Scenario 6: Structured Data Extraction

Task statement 4.5 | Intermediate | Single answer

Correct answer: B - Using custom_id, chunk the 3 oversized documents and resubmit them, and refine the prompt for the 9 validation failures, resubmitting only those.

Why the correct answer is correct: Handle batch failures by resubmitting only failed documents identified by custom_id, with modifications such as chunking oversized documents or refining prompts.

Option-by-option:

A. **[INCORRECT]** Switch all 100 documents to the synchronous API and retry them one at a time.
That gives up batch savings and doesn't address the causes.

B. **[CORRECT]** Using custom_id, chunk the 3 oversized documents and resubmit them, and refine the prompt for the 9 validation failures, resubmitting only those.
Correct. Each failure group gets the fix that matches its cause.

C. **[INCORRECT]** Resubmit all 100 documents unchanged in a new batch and hope that the failures were transient, since the previous run may have hit temporary limits.
The failures have deterministic causes and the 88 successes are already done.

D. **[INCORRECT]** Resubmit only the 3 oversized documents and discard the 9 validation failures as noise.
Discarding validation failures loses data.

Key concept: Batch failure handling (Task 4.5, 4.4).
Exam reasoning: The two groups have different causes.
Production lesson: Diagnose before you resubmit.

Q55 | Domain 4: Prompt Engineering & Structured Output | Scenario 6: Structured Data Extraction

Task statement 4.4 | Challenging | Single answer

Correct answer: D - Extract stated_total and calculated_total, add a conflict_detected boolean when the source itself disagrees, and route flagged records to review.

Why the correct answer is correct: Self-correction flows extract calculated_total alongside stated_total to flag discrepancies, and add conflict_detected booleans for inconsistent source data.

Option-by-option:

A. **[INCORRECT]** Extract only the stated total and trust it, because documents are generally correct and a second total only adds confusion for reviewers downstream.
Inconsistencies would go undetected.

B. **[INCORRECT]** Ask the model to say 'I am confident' or 'I am not confident' at the end of its response.
A confidence statement doesn't identify specific inconsistencies.

C. **[INCORRECT]** Extract only the calculated total from line items and discard the stated total.
Discarding the stated total hides document discrepancies.

D. **[CORRECT]** Extract stated_total and calculated_total, add a conflict_detected boolean when the source itself disagrees, and route flagged records to review.
Correct. It captures both numeric discrepancies and source-level conflicts.

Key concept: Validation flow design (Task 4.4).
Exam reasoning: Both numeric and source-level conflicts need explicit fields.
Production lesson: Design the schema to expose disagreement.

Q56 | Domain 5: Context Management & Reliability | Scenario 6: Structured Data Extraction

Task statement 5.6 | Challenging | Single answer

Correct answer: A - Store each field with its source document, location, excerpt, and date, and annotate conflicts with the competing values and sources instead of silently picking one.

Why the correct answer is correct: Preserve structured claim-source mappings through merging, annotate conflicts with source attribution, and include dates.

Option-by-option:

A. **[CORRECT]** Store each field with its source document, location, excerpt, and date, and annotate conflicts with the competing values and sources instead of silently picking one.
Correct. Provenance and conflict annotation enable review and trust.

B. **[INCORRECT]** Keep the most recently dated value for every field and discard the others, so that the profile stays clean and reviewers see only one value per field, which is also the simplest design.
Silently discarding values hides disagreements.

C. **[INCORRECT]** Have the model write a paragraph summarizing the profile, with sources named in passing.
Prose loses structured traceability.

D. **[INCORRECT]** Add a single note listing all three documents as sources for the whole profile.
A document list doesn't map fields to sources.

Key concept: Provenance and conflict handling (Task 5.6).
Exam reasoning: Reviewers need field-level traceability and visibility of disagreement.
Production lesson: Keep the evidence attached to the claim.

Q57 | Domain 5: Context Management & Reliability | Scenario 6: Structured Data Extraction

Task statement 5.6 | Intermediate | Select 2

Correct answer: A, C - Present the quarterly financial figures as a table. / Present the technical findings as a structured list with their sources.

Why the correct answer is correct: Render different content types appropriately in synthesis outputs: financial data as tables, news as prose, technical findings as structured lists.

Option-by-option:

A. **[CORRECT]** Present the quarterly financial figures as a table.
Correct. Numeric comparisons are clearest in tables.

B. **[INCORRECT]** Present the news developments as a table with one cell per sentence.
News is narrative and suits prose.

C. **[CORRECT]** Present the technical findings as a structured list with their sources.
Correct. Discrete findings suit structured lists.

D. **[INCORRECT]** Render all content as one uniform bulleted list for consistency.
Uniform formatting obscures differences between content types.

E. **[INCORRECT]** Convert the financial figures into a flowing narrative paragraph to make the report more engaging.
Numbers in prose are harder to compare.

Key concept: Content-appropriate rendering (Task 5.6).
Exam reasoning: Each content type has a natural form.
Production lesson: Match the form to the content.

Q58 | Domain 5: Context Management & Reliability | Scenario 6: Structured Data Extraction

Task statement 5.5 | Foundational | Single answer

Correct answer: A - To measure the real error rate among records that the system trusts and to detect new error patterns that confidence scores do not reveal.

Why the correct answer is correct: Stratified random sampling of high-confidence extractions measures ongoing error rates and detects novel error patterns.

Option-by-option:

A. **[CORRECT]** To measure the real error rate among records that the system trusts and to detect new error patterns that confidence scores do not reveal.
Correct. Stratified random sampling supports measurement and discovery.

B. **[INCORRECT]** To reduce the number of low-confidence records that need review by retraining the reviewers, which management prefers.
Sampling is for measurement, not workload.

C. **[INCORRECT]** To confirm that the confidence threshold can be raised without checking any real outcomes.
Confidence scores don't need to be 1.0 for acceptance.

D. **[INCORRECT]** To avoid having to calibrate confidence thresholds on a labeled set.
Calibration is still needed. Sampling complements it.

Key concept: Sampling rationale (Task 5.5).
Exam reasoning: The accepted population is where unseen errors live.
Production lesson: Audit what you auto-accept.

Q59 | Domain 5: Context Management & Reliability | Scenario 6: Structured Data Extraction

Task statement 5.3 | Intermediate | Single answer

Correct answer: A - Commit the 92 valid records, report the 8 failures with their reasons for correction or review, and keep the rest of the run's results.

Why the correct answer is correct: Avoid both silent suppression and terminating entire workflows on single failures; return partial results with structured failure details.

Option-by-option:

A. **[CORRECT]** Commit the 92 valid records, report the 8 failures with their reasons for correction or review, and keep the rest of the run's results.
Correct. Partial success with explicit failure reporting preserves value and visibility.

B. **[INCORRECT]** Commit the 92 and silently drop the 8 so the run appears to succeed.
Silent drops lose data and hide problems.

C. **[INCORRECT]** Keep rolling back all 100 so the downstream system is always consistent with a single run.
This discards 92 good records for the sake of 8.

D. **[INCORRECT]** Retry the whole run of 100 repeatedly until all pass.
Records with deterministic problems will never pass.

Key concept: Partial results in downstream integration (Task 5.3).
Exam reasoning: Only 8% failed, and the reasons are specific.
Production lesson: Fail small and loudly.

Q60 | Domain 5: Context Management & Reliability | Scenario 6: Structured Data Extraction

Task statement 5.5 | Challenging | Single answer, secondary: D4

Correct answer: B - Route it to human review because contradictory source data warrants review regardless of the confidence score.

Why the correct answer is correct: Extractions with ambiguous or contradictory source documents should be routed to human review, in addition to low model confidence.

Option-by-option:

A. **[INCORRECT]** Auto-accept it, since 0.97 exceeds the threshold and thresholds are the agreed rule.
The rule fails to account for the conflict flag.

B. **[CORRECT]** Route it to human review because contradictory source data warrants review regardless of the confidence score.
Correct. Confidence does not resolve a conflict in the source.

C. **[INCORRECT]** Reject it permanently because documents with conflicting totals are invalid.
A human can often resolve the conflict, so rejection discards recoverable data.

D. **[INCORRECT]** Auto-accept it but lower the confidence to 0.90 so that dashboards show uncertainty.
Adjusting a number doesn't resolve the conflict.

Key concept: Routing on conflicts (Task 5.5, 4.4).
Exam reasoning: The conflict flag overrides a high confidence number.
Production lesson: Confidence is not truth.

---

7. Domain-level performance review guide

Use your tracker. A domain below roughly 70% on a paper deserves a focused review before the next paper. Look at the task statements and concepts listed for each domain and re-read the matching sections of the official Exam Guide.

Domain 1 - Agentic Architecture & Orchestration (27% official weight; 16 questions here)

Task	Topic	Questions
1.1	Agentic loops for autonomous task execution	14, 15, 18
1.2	Coordinator-subagent orchestration	17, 19
1.3	Subagent invocation, context passing, spawning	16, 20
1.5	Agent SDK hooks: interception and normalization	44, 45, 46
1.6	Task decomposition strategies	30, 31, 32
1.7	Session state, resumption, and forking	1, 2, 21

Lessons to carry forward: Carry forward conclusions, not stale evidence. Fork to compare, resume to continue. Know the two key stop reasons cold. Guards should fail visibly. A subagent can only use what it is handed. Give each worker its own lane.

Domain 2 - Tool Design & MCP Integration (18% official weight; 11 questions here)

Task	Topic	Questions
2.1	Tool interfaces: descriptions and boundaries	25, 26, 50
2.2	Structured error responses for MCP tools	22, 33, 49
2.3	Tool distribution and tool_choice	24, 47, 48
2.4	MCP server integration (Claude Code and agents)	23, 34

Lessons to carry forward: Help the agent choose a different path. Sell the tool to the model in its description. Tool access defines behavior. Describe usage, not marketing. Make their differences visible. Make the error say 'stop' and 'try this instead'.

Domain 3 - Claude Code Configuration & Workflows (20% official weight; 12 questions here)

Task	Topic	Questions
3.1	CLAUDE.md hierarchy, scoping, modularity	3
3.2	Custom slash commands and skills	4, 8
3.3	Path-specific rules	6, 9
3.4	Plan mode vs direct execution	7
3.5	Iterative refinement techniques	5, 10
3.6	Claude Code in CI/CD pipelines	35, 36, 37, 38

Lessons to carry forward: Choose scope by who and which files. Know where each option lives. Match the technique to the failure. Rules follow file patterns. Small and clear: just do it. Learn what each frontmatter field actually does.

Domain 4 - Prompt Engineering & Structured Output (20% official weight; 12 questions here)

Task	Topic	Questions
4.1	Explicit criteria; reducing false positives	41
4.2	Few-shot prompting	11, 12
4.3	Structured output via tool use and JSON schemas	53
4.4	Validation, retry, feedback loops	51, 55
4.5	Batch processing (Message Batches API)	52, 54
4.6	Multi-instance and multi-pass review	27, 39, 40, 42

Lessons to carry forward: Show how to decide, not just what to do. Examples are instructions. Curate them. Independence and focus beat instruction. Independent means uninformed of the author's logic. Fix the cause of noise, not the symptom. Define in or out, not how sure.

Domain 5 - Context Management & Reliability (15% official weight; 9 questions here)

Task	Topic	Questions
5.3	Error propagation across multi-agent systems	28, 29, 59
5.4	Context in large codebase exploration	13
5.5	Human review workflows; confidence calibration	43, 58, 60
5.6	Provenance and uncertainty in multi-source synthesis	56, 57

Lessons to carry forward: Compact context, keep durable notes elsewhere. Show the shape of the evidence. Report absence honestly and with its method. Break the average apart. Keep the evidence attached to the claim. Match the form to the content.

8. Final revision checklist

Tick each item when you can explain it without looking at the key. Items are drawn from the concepts tested in this paper.

Done	Concept	First seen
[ ]	Fresh session with injected summary	Q1
[ ]	fork_session use cases	Q2
[ ]	Configuration scope mapping	Q3
[ ]	Skill frontmatter	Q4
[ ]	Iterative refinement techniques	Q5
[ ]	Path-scoped rules	Q6
[ ]	Mode selection	Q7
[ ]	Skills and commands	Q8
[ ]	Conditional rule loading	Q9
[ ]	Test-driven iteration	Q10
[ ]	Few-shot for ambiguous decisions	Q11
[ ]	Example quality	Q12
[ ]	/compact	Q13
[ ]	Loop control flow	Q14
[ ]	Loop termination with backstops	Q15
[ ]	Context packets	Q16
[ ]	Scope partitioning	Q17
[ ]	Hook outcomes in the agentic loop	Q18
[ ]	Parallel versus sequential orchestration	Q19
[ ]	No shared memory between invocations	Q20
[ ]	Stale results versus fresh start	Q21
[ ]	Permission errors and alternatives	Q22
[ ]	MCP tool descriptions	Q23
[ ]	Tool distribution across agents	Q24
[ ]	Description content	Q25
[ ]	Eliminating overlap	Q26
[ ]	Multi-instance and multi-pass reviews	Q27
[ ]	Coverage annotations	Q28
[ ]	Valid empty results	Q29
[ ]	Prompt chaining for test generation	Q30
[ ]	Cross-file integration review	Q31
[ ]	Adaptive decomposition	Q32
[ ]	Permission errors in CI	Q33
[ ]	MCP credentials	Q34
[ ]	CI invocation	Q35
[ ]	Evaluation of CI review quality	Q36
[ ]	CI versus delivery versus deployment	Q37
[ ]	Least privilege in CI	Q38
[ ]	Review independence	Q39
[ ]	Precision versus recall trade-offs	Q40

Done	Concept	First seen
[ ]	Explicit criteria	Q41
[ ]	Multi-pass architecture	Q42
[ ]	Segmented accuracy	Q43
[ ]	Hooks for guaranteed behavior	Q44
[ ]	PostToolUse normalization	Q45
[ ]	Deterministic guardrails for batch processing	Q46
[ ]	Scoped tool access	Q47
[ ]	Least privilege for side effects	Q48
[ ]	Validation errors	Q49
[ ]	Precise tool contracts	Q50
[ ]	Bounded retries	Q51
[ ]	Batch limitations	Q52
[ ]	Monetary schema design	Q53
[ ]	Batch failure handling	Q54
[ ]	Validation flow design	Q55
[ ]	Provenance and conflict handling	Q56
[ ]	Content-appropriate rendering	Q57
[ ]	Sampling rationale	Q58
[ ]	Partial results in downstream integration	Q59
[ ]	Routing on conflicts	Q60

Before the next paper: classify each wrong answer as (a) knowledge gap, (b) misread of the stem or qualifier, or (c) a plausible-distractor trap. See CCAF_Final_Revision_Guide.pdf for the method.